

Advanced Machine Learning



Bogdan Alexe,

bogdan.alex@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2025-2026

Administrative

- first seminar class:
 - today, 10-12, room 209 + next week the same seminar
 - next Monday, 18-20, room 209 - the same seminar
- I've uploaded the exercises for the first seminar (taken from the course book)

Today's lecture: Overview

- A Formal Model – The Statistical learning framework
- Empirical Risk Minimization
- Probably Approximately Correct learning

Particular learning scenario – papaya tasting


**You've just arrived in some
small Pacific island**



**You soon find that papayas are a
significant ingredient in the local diet**



**Based on previous experience with other
fruits, you decide to use two features**



Formal model for statistical learning

Example


- **domain set, $\mathcal{X}$** : this is the set of objects that we may wish to label.
 - instance space
 - elements of $\mathcal{X}$ are called instances, represented by *features*
 - **label set, $\mathcal{Y}$** : the set of possible labels.
 - usually $\{0,1\}$ or $\{-1, +1\}$
 - **training data S**
 - the learner's input
 - supervised batch learning scenario
 - finite sequence of pairs S of labeled domain points: $S = ((x_1, y_1) \dots (x_m, y_m)) \in (\mathcal{X} \times \mathcal{Y})^m$
- $\mathcal{X} = \mathbf{R}^2$ representing color and shape of papays.
 - $\mathcal{Y} = \{0, 1\}$ representing “tasty” (1) or “non-tasty”(0)
 - S = set of already tasted papayas and their color, softness and tastiness (label)

Formal model for statistical learning

- a prediction rule, $h : \mathcal{X} \rightarrow \mathcal{Y}$
 - the learner's output;
 - used to label future examples;
 - called a *predictor*, a *hypothesis*, or a *classifier*;
 - $h = A(S)$: the hypothesis h is returned by a learning algorithm A based on the training sequence S
 - *goal of the learner*: h should be *correct on future examples*

Example

- prediction rule for tastiness
- $h(x) = 1$ if $x = [\text{color}, \text{shape}]$ within the inner rectangle,
- $h(x) = 0$ if $x = [\text{color}, \text{shape}]$ outside the inner rectangle.



“Correct on future examples”

- let f be the correct classifier, then we should find h such that $h \approx f$
- one way: define the error of h w.r.t. f to be the probability that it does not predict the correct label on a random data point x generated by the underlying (unknown) probability distribution $\mathcal{D}$ over $\mathcal{X}$:

$$L_{\mathcal{D},f}(h) = \mathbb{P}_{x \sim \mathcal{D}}[h(x) \neq f(x)]$$

- more formally, $\mathcal{D}$ is a probability distribution over $\mathcal{X}$, that is, for a given $A \subset \mathcal{X}$, the value of $\mathcal{D}(A)$ is the probability to observe some $x \in A$. Then:

$$L_{\mathcal{D},f}(h) \stackrel{\text{def}}{=} \mathbb{P}_{x \sim \mathcal{D}}[h(x) \neq f(x)] \stackrel{\text{def}}{=} \mathcal{D}(\{x \in \mathcal{X} : h(x) \neq f(x)\})$$

- can we find h s.t. $L_{\mathcal{D},f}(h)$ is small ?
 - $L_{\mathcal{D},f}(h)$ is called **generalization error**, the **true error** of h , the **real risk** of h
 - L - loss of the learner

Data-generation model

- we must assume some relation between the training data S and $\mathcal{D}, f$
- simple data generation model:
 - **Independently Identically Distributed (i.i.d.)**: Each x_i is sampled independently according to $\mathcal{D}$.
 - we do not assume that the learner knows anything about $\mathcal{D}$.
 - **Realizability**: assume that there is a “correct” labeling function, $f: \mathcal{X} \rightarrow \mathcal{Y}$ and that for every $i \in 1, 2, \dots, m$, $y_i = f(x_i)$
 - f is a deterministic labeling – strong assumption
 - relax this assumption later (make f probabilistic labeling)
 - we do not assume that the learner knows anything about f .
- each pair in the training data S is generated by first sampling i.i.d. a point x_i according to $\mathcal{D}$ and then labeling it by f .

Empirical Risk Minimization

The ERM learning paradigm

- $h_S = A(S)$: the hypothesis h_S is returned by a learning algorithm A based on the training sequence S , sampled i.i.d from an unknown distribution $\mathcal{D}$ and labeled by some target function f
- the true error $L_{\mathcal{D},f}(h)$ is unknown to the learner as he doesn't know $\mathcal{D}$ and f
- the learner has access to the *training error* = *empirical error* = *empirical risk*:

$$L_S(h) \stackrel{\text{def}}{=} \frac{|\{i \in [m] : h(x_i) \neq y_i\}|}{m},$$


- search for a solution that works well on the training data – minimize $L_S(h)$
- *Empirical Risk Minimization* (ERM) = learning paradigm that returns a predictor h that minimizes $L_S(h)$

ERM might overfit

- simple rule for finding h with small empirical error:


$$h_S(x) = \begin{cases} y_i & \text{if } \exists i \in [m] \text{ s.t. } x_i = x \\ 0 & \text{otherwise.} \end{cases}$$


- consider:
 - $\mathcal{D}$ uniform over the gray rectangle
 - all instances in the inner rectangle are labeled 1 (**blue points**)
 - all other instances in the gray rectangle but outside the inner one are labeled 0 (**red points**)
 - area of the gray square is 2, area of the inner rectangle is 1
- we obtain that $L_S(h_S) = 0$, but $L_{\mathcal{D},f}(h_S) = 1/2$ (h predicts the label 1 on a finite number of instances)
- *overfitting*: small training error, large true error
- while the above predictor h_S seems to be very unnatural, it can be described as a thresholded polynomial (seminar exercise)



ERM with Inductive Bias

- to guard against overfitting we introduce some prior knowledge (inductive bias)
- example for the papaya prediction tasting problem: there exists a good prediction rule that is some axis aligned rectangle
- restricted search space: set of all rectangles
- hypothesis class = $\mathcal{H} \subset \mathcal{Y}^{\mathcal{X}} = \{g : \mathcal{X} \rightarrow \mathcal{Y}\}$
- revised ERM rule: apply the ERM learning paradigm over $\mathcal{H}$
 - for the training sample S , the $\text{ERM}_{\mathcal{H}}$ learner chooses a predictor $h \in \mathcal{H}$ with the lowest possible error over S :
$$\text{ERM}_{\mathcal{H}}(S) \in \underset{h \in \mathcal{H}}{\text{argmin}} L_S(h),$$
- over which hypothesis classes $\text{ERM}_{\mathcal{H}}$ will not result in overfitting?
 - what characterizes a good hypothesis class?






PAC learning

(Probably Approximately Correct learning)

PAC learning

Leslie Valiant, Turing award 2010

*For transformative contributions to the theory of computation, including **the theory of probably approximately correct (PAC) learning**, the complexity of enumeration and of algebraic computation, and the theory of parallel and distributed computing.*



PAC learning


Leslie Valiant, Turing award 2010

*For transformative contributions to the theory of computation, including **the theory of probably approximately correct (PAC) learning**, the complexity of enumeration and of algebraic computation, and the theory of parallel and distributed computing.*

PAC learning (Probably Approximately Correct learning) is a theoretical framework introduced by Leslie Valiant in 1984 that formalizes what it means for a learning algorithm to succeed in a ***distribution free setting***: the goal is to design an algorithm that, given a set of training examples drawn from an unknown distribution, will produce a hypothesis that is "probably approximately correct"—that is, its error is below a pre-specified threshold (ϵ) – with high probability (at least $1 - \delta$).

Can only be *Approximately* correct

- **claim:** we can't hope to find h_S s.t. $L_{(\mathcal{D},f)}(h_S) = 0$
- **proof:**
 - for every $\varepsilon \in (0, 1)$ take $\mathcal{X} = \{x_1, x_2\}$ and $\mathcal{D}(\{x_1\}) = 1 - \varepsilon$, $\mathcal{D}(\{x_2\}) = \varepsilon$
 - the probability not to see x_2 at all among m i.i.d. examples from S is $(1 - \varepsilon)^m \approx e^{-\varepsilon m}$



Can only be *Approximately* correct

- **claim:** we can't hope to find h_S s.t. $L_{(\mathcal{D},f)}(h_S) = 0$
- **proof:**
 - for every $\varepsilon \in (0, 1)$ take $\mathcal{X} = \{x_1, x_2\}$ and $\mathcal{D}(\{x_1\}) = 1 - \varepsilon$, $\mathcal{D}(\{x_2\}) = \varepsilon$
 - the probability not to see x_2 at all among m i.i.d. examples from S is $(1 - \varepsilon)^m \approx e^{-\varepsilon m}$
- so, if $\varepsilon \ll 1/m$ we're likely not to see x_2 at all, but then we can't know its label
- **relaxation:** we'd be happy with $L_{(\mathcal{D},f)}(h_S) \leq \varepsilon$, where ε is a parameter user-specified
- ε is the accuracy parameter
- ε measures the quality of our prediction

Can only be *Probably* correct

- recall that the input to the learner is randomly generated
- there's always a (very small) chance to see the same example again and again
 - all the papayas we have happened to taste were not tasty,
 - we will come up with the classifier that assigns label 0 = Not Tasty to every future sample
 - will lead to large true error
- **claim:** no algorithm can guarantee $L_{(\mathcal{D},f)}(h_S) \leq \varepsilon$ for sure
 - address the probability to sample a training set S for which $L_{(\mathcal{D},f)}(h_S) \leq \varepsilon$
- **relaxation:** we allow the algorithm to fail with probability δ , where $\delta \in (0, 1)$ is user-specified
- here, the probability δ is over the random choice of examples
- $1 - \delta$ is the **confidence parameter**
- δ measures the probability of getting a non-representative sample

$$P_{S \sim D^m} (L_{f,D}(h_S) \leq \varepsilon) \geq 1 - \delta \Leftrightarrow P_{S \sim D^m} (L_{f,D}(h_S) > \varepsilon) < \delta$$

Probably **A**pproximately **C**orrect (**PAC**) learning

- the learner doesn't know $\mathcal{D}$ and f
- the learner receives accuracy parameter ε and confidence parameter δ
- the learner can ask for training data, S , containing $m(\varepsilon, \delta)$ examples (that is, the number of examples can depend on the value of ε and δ , but it can't depend on $\mathcal{D}$ or f).
- learner should output a hypothesis h_S s.t. with probability of at least $1 - \delta$ it holds that $L_{\mathcal{D},f}(h_S) \leq \varepsilon$
- that is, the learner should be **P**robably (with probability at least $1 - \delta$) **A**pproximately (up to accuracy ε) **C**orrect

$$P_{S \sim D^m}(L_{f,D}(h_S) \leq \varepsilon) \geq 1 - \delta \Leftrightarrow P_{S \sim D^m}(L_{f,D}(h_S) > \varepsilon) < \delta$$

Learning finite classes

- give more knowledge to the learner: the target f comes from some **hypothesis class**, $\mathcal{H} \subset \mathcal{Y}^{\mathcal{X}}$
- the learner knows $\mathcal{H}$
- assume that $\mathcal{H}$ is a finite hypothesis class
 - example: $\mathcal{H}$ is the set of conjunctions with d variable (binary features), $|\mathcal{H}| = 3^d + 1$
- use the **ERM** learning rule:
 - input: $\mathcal{H}$ and $S = ((x_1, y_1), \dots, (x_m, y_m))$
 - output: $h \in \mathcal{H}$, h is an ERM hypothesis
- **Empirical Risk Minimization (ERM)**
 - input: training sequence $S = ((x_1, y_1), \dots, (x_m, y_m))$
 - define the empirical risk $L_S(h)$:
$$L_S(h) = \frac{|\{i \in [m] : h(x_i) \neq y_i\}|}{m},$$
 - output: any $h \in \mathcal{H}$ that minimizes $L_S(h)$

The Realizability Assumption

- assume that there exists $h^* \in \mathcal{H}$ such that $L_{\mathcal{D},f}(h^*) = 0$.
 - implies that with probability 1 over random samples, S , where the instances of S are sampled according to $\mathcal{D}$ and are labeled by f , we have $L_S(h^*) = 0$.
 - implies that for every ERM hypothesis h_S we have that $L_S(h_S) = 0$
 - h_S has minimum empirical risk, but we already know that h^* has empirical risk equal to 0
- also known as the consistency assumption
- strong assumption
 - relax the assumption later

Learning finite classes

Theorem:

Let $\mathcal{H}$ be a finite hypothesis class. Let $\delta \in (0, 1)$ and $\epsilon > 0$ and let m be an integer that satisfies:

$$m \geq \frac{\log(|\mathcal{H}|/\delta)}{\epsilon}$$

Then, for any labeling function f , and for any distribution $\mathcal{D}$, for which the realizability assumption holds (that is, for some $h^* \in \mathcal{H}$, $L_{\mathcal{D},f}(h^*) = 0$) with probability of least $1 - \delta$ over the choice of an i.i.d, sample S of size m , we have that for every ERM hypothesis h_S it holds that: $L_{(\mathcal{D},f)}(h_S) \leq \epsilon$.

$$P_{S \sim \mathcal{D}^m}(L_{f,D}(h_S) \leq \epsilon) \geq 1 - \delta$$

*The theorem basically says that for a sufficient large number of training examples m the $ERM_{\mathcal{H}}$ rule over a finite hypothesis class is **Probably Approximately Correct***

Learning finite classes

*The theorem basically says that for a sufficient large number of training examples m the $ERM_{\mathcal{H}}$ rule over a finite hypothesis class is **Probably Approximately Correct***

$$P_{S \sim D^m} (L_{f,D}(h_S) \leq \epsilon) \geq 1 - \delta$$

Idea of the proof

- a bad predictor h_b has $L_{(D,f)}(h_b) > \epsilon$
- h_b can be output by the $ERM_{\mathcal{H}}$ learning paradigm if has zero empirical error: $L_S(h_b) = 0$
- this can happen if h_b labels correctly all the m training examples from S i.i.d from $\mathcal{D}$
- given a random example from $\mathcal{D}$, h_b has $< 1-\epsilon$ probability to label it correctly
- h_b labels correctly all the m training examples from S with probability $< (1-\epsilon)^m \leq e^{-\epsilon m}$
- there are at most $|\mathcal{H}|$ bad hypothesis, so consider $|\mathcal{H}| \times e^{-\epsilon m} \leq \delta$, so take $m \geq \frac{\log(|\mathcal{H}|/\delta)}{\epsilon}$

PAC learnability of a class $\mathcal{H}$ -definition

A hypothesis class $\mathcal{H}$ is called **PAC learnable** if there exists a function $m_{\mathcal{H}} : (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm A with the following property:

- for every $\varepsilon > 0$ (*accuracy* $\rightarrow$ “approximately correct”)
- for every $\delta > 0$ (*confidence* $\rightarrow$ “probably”)
- for every labeling $f \in \mathcal{H}$ (*realizability case*)
- for every distribution $\mathcal{D}$ over $\mathcal{X}$

when we run the learning algorithm A on a training set S , consisting of $m \geq m_{\mathcal{H}}(\varepsilon, \delta)$ examples sampled i.i.d. from $\mathcal{D}$ and labeled by f the algorithm A returns a hypothesis $h_S \in \mathcal{H}$ such that, with probability at least $1-\delta$ (over the choice of examples), $L_{\mathcal{D},f}(h_S) \leq \varepsilon$.

PAC learnability of a class $\mathcal{H}$ -definition

A hypothesis class $\mathcal{H}$ is called **PAC learnable** if there exists a function $m_{\mathcal{H}} : (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm A with the following property:

- for every $\varepsilon > 0$ (*accuracy* $\rightarrow$ “approximately correct”)
- for every $\delta > 0$ (*confidence* $\rightarrow$ “probably”)
- for every labeling $f \in \mathcal{H}$ (*realizability case*)
- for every distribution $\mathcal{D}$ over $\mathcal{X}$

when we run the learning algorithm A on a training set S , consisting of $m \geq m_{\mathcal{H}}(\varepsilon, \delta)$ examples sampled i.i.d. from $\mathcal{D}$ and labeled by f the algorithm A returns a hypothesis $h_S \in \mathcal{H}$ such that, with probability at least $1-\delta$ (over the choice of examples), $L_{\mathcal{D},f}(h_S) \leq \varepsilon$.

$$P_{S \sim \mathcal{D}^m} (L_{f,D}(h_S) \leq \varepsilon) \geq 1 - \delta$$

- $h_S = A(S)$
- the function $m_{\mathcal{H}} : (0,1)^2 \rightarrow \mathbb{N}$ is called sample complexity of learning $\mathcal{H}$
- $m_{\mathcal{H}}(\varepsilon, \delta)$ – the minimum number of examples required to guarantee a PAC solution

PAC learnability of a class $\mathcal{H}$ -definition

A hypothesis class $\mathcal{H}$ is called **PAC learnable** if there exists a function $m_{\mathcal{H}} : (0,1)^2 \rightarrow \mathbb{N}$ and a learning algorithm A with the following property:

- for every $\varepsilon > 0$ (*accuracy* $\rightarrow$ “approximately correct”)
- for every $\delta > 0$ (*confidence* $\rightarrow$ “probably”)
- for every labeling $f \in \mathcal{H}$ (*realizability case*)
- for every distribution $\mathcal{D}$ over $\mathcal{X}$

when we run the learning algorithm A on a training set S , consisting of $m \geq m_{\mathcal{H}}(\varepsilon, \delta)$ examples sampled i.i.d. from $\mathcal{D}$ and labeled by f the algorithm A returns a hypothesis $h_S \in \mathcal{H}$ such that, with probability at least $1-\delta$ (over the choice of examples), $L_{\mathcal{D},f}(h_S) \leq \varepsilon$.

$$P_{S \sim \mathcal{D}^m}(L_{f,D}(h_S) \leq \varepsilon) \geq 1 - \delta \Leftrightarrow P_{S \sim \mathcal{D}^m}(L_{f,D}(h_S) > \varepsilon) < \delta$$

Sample complexity


- the function $m_{\mathcal{H}}: (0,1)^2 \rightarrow \mathbb{N}$ is called sampled complexity of learning $\mathcal{H}$
- $m_{\mathcal{H}}(\epsilon, \delta)$ – the minimum number of examples required to guarantee a PAC solution
- depends on:
 - accuracy ϵ
 - confidence δ
 - properties of $\mathcal{H}$
- different than *time complexity* (discuss it in the following lectures)

Every finite hypothesis class H is PAC learnable with sample complexity

$$m_{\mathcal{H}}(\epsilon, \delta) \leq \left\lceil \frac{\log(|\mathcal{H}|/\delta)}{\epsilon} \right\rceil$$

Axis-aligned rectangles

- $\mathcal{X} = \mathbb{R}^2$ points in the plane
- $\mathcal{H}$ = set of all axis-aligned rectangle lying in $\mathbb{R}^2$
- each concept $h \in \mathcal{H}$ is an indicator function of a rectangle
- the learning problem consists of determining with small error a target axis-aligned rectangle using the labeled training sample




Target concept R and possible hypothesis R' . Circles represent training instances. A blue circle is a point labeled with 1, since it falls within the rectangle R . Others are red and labeled with 0.

Axis-aligned rectangles

- $\mathcal{X} = \mathbb{R}^2$ points in the plane
- $\mathcal{H}$ = set of all axis-aligned rectangle lying in $\mathbb{R}^2$
- $|\mathcal{H}| = \infty$
- still $\mathcal{H}$ is PAC-learnable with sample complexity in the order of:

$$m \geq \left\lceil \frac{4}{\varepsilon} \log\left(\frac{1}{\delta}\right) \right\rceil$$

- simple algorithm: take the tightest rectangle enclosing all the positive examples (or take the largest rectangle not including negative samples)
- discuss this example in seminar



Next time

3	A Formal Learning Model	43
3.1	PAC Learning	43
3.2	A More General Learning Model	44
3.2.1	Releasing the Realizability Assumption – Agnostic PAC Learning	45
3.2.2	The Scope of Learning Problems Modeled	47
3.3	Summary	49
3.4	Bibliographic Remarks	50
3.5	Exercises	50