

# 4. Debugging

# Debugging software can be painful



<https://quotefancy.com/quote/3753413/IEEE-Computer-Society-Software-and-cathedrals-are-much-the-same-first-we-build-them-then>

# Machine Learning debugging can be even more tedious



Laura Ruis  
@LauraRuis

Follow

...

when you debug your NN for 2 days until you realise the lr is 1e3 instead of 1e-3 😊 cool

5:21 PM · Feb 9, 2023 · 91.3K Views

---

37

18

905

12

↑

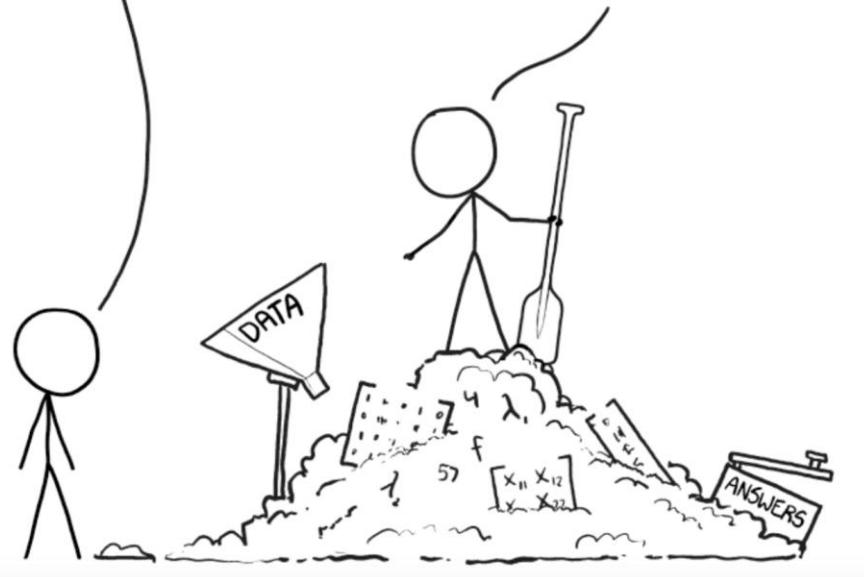
B

THIS IS YOUR MACHINE LEARNING SYSTEM?

YUP! YOU POUR THE DATA INTO THIS BIG  
PILE OF LINEAR ALGEBRA, THEN COLLECT  
THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL  
THEY START LOOKING RIGHT.



<https://xkcd.com/1838/>

Machine Learning development has the reputation of a *trial-and-error, messy process*

If we are principled about it, we can make  
**debugging ML projects more predictable**

# Software Engineering

- **more predictable:** logic is defined by code

# Machine Learning

- **less predictable:** ‘logic’ is defined by code + data + training dynamics

# Software Engineering

- **more predictable:** logic is defined by code
- **fewer sources of errors:** typically logic bugs or integration issues

# Machine Learning

- **less predictable:** ‘logic’ is defined by code + data + training dynamics
- **more sources of error:** data quality, wrong architecture or hyperparameters

# Software Engineering

- **more predictable:** logic is defined by code
- **fewer sources of errors:** typically logic bugs or integration issues
- **explainability:** easier to trace back why system behaved a certain way

# Machine Learning

- **less predictable:** ‘logic’ is defined by code + data + training dynamics
- **more sources of error:** data quality, wrong architecture or hyperparameters
- **explainability:** ML systems are more opaque => need to look at training dynamics, activations, feature importance etc.

# Plan for today

- learn how to tackle the most common bugs in ML lifecycle
  - **syntactic bugs**
  - **functional bugs**
- inspect learning curves:
  - look for signs of good fit/overfitting/underfitting
  - recognize potential ***hyperparameters misconfigurations***

hopefully, by the end of this lecture you will be more confident in when debugging ML projects 

# Machine Learning bugs

syntactic

functional

```
-----  
ValueError                                Traceback (most recent call last)  
<ipython-input-4-59c95019453c> in <module>()  
      1 model = Sequential()  
----> 2 model.add(Conv1D(filters=64, kernel_size=3, activation='relu', input_shape=(50,)))  
      3 model.add(Conv1D(filters=64, kernel_size=3, activation='relu'))  
      4 model.add(Dropout(0.5))  
      5 model.add(MaxPooling1D(pool_size=2))  
  
-/anaconda3/lib/python3.6/site-packages/keras/engine/sequential.py in add(self, layer)  
    163         # and create the node connecting the current layer  
    164             # to the input layer we just created.  
--> 165             layer(x)  
    166             set_inputs = True  
    167         else:  
  
-/anaconda3/lib/python3.6/site-packages/keras/engine/base_layer.py in __call__(self, inputs, **kwargs)  
    412     # Raise exceptions in case the input is not compatible  
    413     # with the input_spec specified in the layer constructor.  
--> 414     self.assert_input_compatibility(inputs)  
    415  
    416     # Collect input shapes to build layer.  
  
-/anaconda3/lib/python3.6/site-packages/keras/engine/base_layer.py in assert_input_compatibility(self, inputs)  
    309         self.name + ': expected ndim=' +  
    310         str(spec.ndim) + ', found ndim=' +  
--> 311         str(K.ndim(x)))  
    312     if spec.max_ndim is not None:  
    313         ndim = K.ndim(x)  
  
ValueError: Input 0 is incompatible with layer conv1d_1: expected ndim=3, found ndim=2
```



# Wrong tensor shape

 Discussion by  Henk Poley

**tensor 'token\_embd.weight' has wrong shape**

TheBloke/CodeLlama-7B-Python-GGUF

- one of the most common Deep Learning errors
- usually verbose, but can lead to silent bugs

# Wrong tensor shape - wrong input shape

```
class SimpleMLP(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(10, 20)
        self.act1 = nn.ReLU()
        self.fc2 = nn.Linear(20, 2)

    def forward(self, x):
        h = self.act1(self.fc1(x))
        o = self.fc2(h)
        return o
```

```
model = SimpleMLP()
x = torch.rand(3, 11)
model(x)
```

```
8     def forward(self, x):
----> 9         h = self.act1(self.fc1(x))
10        o = self.fc2(h)
11        return o
```

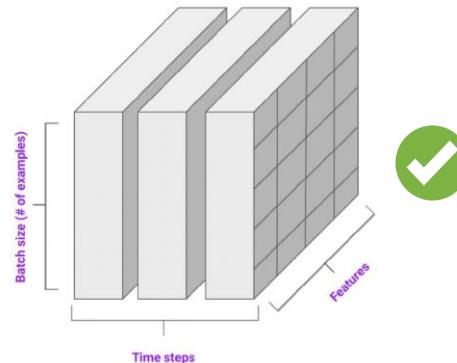
- a lot of the times, the *input data* has the wrong shape

RuntimeError: mat1 and mat2 shapes cannot be multiplied (3x11 and 10x20)

# Wrong tensor shape - wrong input shape

typical scenarios:

- *images*:
  - input Tensor has shape  
**batch x W x H x channels**  
instead of  
**batch x channels x W x H**
  - inputs are greyscale and  
have **batch x W x H**
- *sequential data*: input Tensor has  
shape **batch x H x timesteps**  
instead of **batch x timesteps x H**



# Wrong tensor shape - wrong input shape

scenario: last batch in a dataset may have 1st dimension != **batch\_size**:

- why? let's say **batch\_size=32**
- **DataLoader** loads 32 elements until it exhausts all the dataset elements
- if dataset has 3216 elements => **last batch** has 16 elements

# Wrong tensor shape - wrong input shape

scenario: last batch in a dataset may have 1st dimension != **batch\_size**:

- if code assumes output size is **batch\_size**, shape errors may appear:

```
# resize images: 32x3x32x32 => 32x(3x32x32)
print("[train_epoch] batch_img before resize = ", batch_img.size())
batch_img = batch_img.view(32, -1)
```

```
[train_epoch] batch_img before resize =  torch.Size([16, 3, 32, 32])
[train_epoch] batch_img after resize =  torch.Size([32, 1536])
```

new input size (1536) no  
longer matches layer size  
( $3 \times 3 \times 32 = 3072$ )

```
RuntimeError: mat1 and mat2 shapes cannot be multiplied (32x1536 and 3072x256)
```

# Wrong tensor shape - wrong label shape

labels are one-hot but loss function expects scalar **[0, num\_classes-1]**

Label Encoding			One Hot Encoding			
Food Name	Categorical #	Calories	Apple	Chicken	Broccoli	Calories
Apple	1	95	1	0	0	95
Chicken	2	231	0	1	0	231
Broccoli	3	50	0	0	1	50

source: <https://www.thetalkingmachines.com/article/beginners-guide-debugging-tensorflow-models>

# Wrong tensor shape - wrong output shape

output doesn't conform to *shape expected by loss function*

The *input* is expected to contain raw, unnormalized scores for each class. *input* has to be a Tensor of size either  $(\text{minibatch}, C)$  or  $(\text{minibatch}, C, d_1, d_2, \dots, d_K)$  with  $K \geq 1$  for the  $K$ -dimensional case. The latter is useful

documentation snippet from [CrossEntropyLoss](#) in Pytorch

# Wrong tensor shape - wrong output shape

typical scenarios:

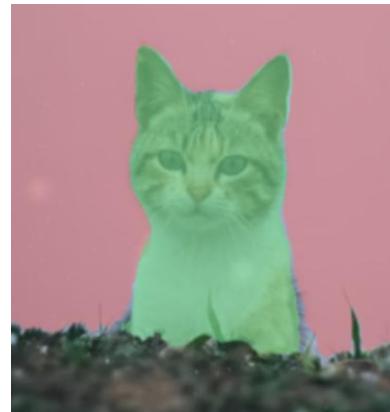
- sequential data:
  - if each timestep is classified into C classes
  - output for CrossEntropyLoss should be **batch x C x timesteps**, not **batch x timesteps x C**
- semantic segmentation:
  - C classes
  - output for CrossEntropyLoss should be **?**



# Wrong tensor shape - wrong output shape

typical scenarios:

- sequential data:
  - if each timestep is classified into C classes
  - output for CrossEntropyLoss should be **batch x C x timesteps**, not **batch x timesteps x C**
- semantic segmentation:
  - C classes
  - output for CrossEntropyLoss should be **batch x C x H x W**



# Wrong tensor shape - intermediate ops

- during feedforward, tensors are manipulated using different ops (**slice**, **concat**, **max**, **reshape** etc.)
- common scenario:
  - calling functions with the **wrong arguments** *alters tensor shapes unexpectedly*
  - the resulting **tensor shape no longer matches with the next layer**

# Wrong tensor shape - intermediate ops



# Wrong tensor shape - intermediate ops

scenario: **concatenating** two tensors on the *wrong dimension*

```
BATCH_SIZE = 16
# batch_size x 50
h1 = torch.rand(BATCH_SIZE, 50)
h2 = torch.rand(BATCH_SIZE, 50)

h_cat = torch.cat((h1, h2))
print(h_cat.size())

torch.Size([32, 50])
```



cat function concatenates on **dim=0** by default!

# Wrong tensor shape - intermediate ops

scenario: **concatenating** two tensors on the *wrong dimension*

good practice: specify dimensions explicitly whenever possible (**dim** argument)

```
BATCH_SIZE = 16
# batch_size x 50
h1 = torch.rand(BATCH_SIZE, 50)
h2 = torch.rand(BATCH_SIZE, 50)

h_cat = torch.cat((h1, h2))
print(h_cat.size())

torch.Size([32, 50])
```



cat function concatenates on **dim=0** by default!

```
BATCH_SIZE = 16
# batch_size x 50
h1 = torch.rand(BATCH_SIZE, 50)
h2 = torch.rand(BATCH_SIZE, 50)

h_cat = torch.cat((h1, h2), dim=1)
print(h_cat.size())

torch.Size([16, 100])
```



the two inputs are correctly concatenated

# Wrong tensor shape - intermediate ops

scenario:

- tensor has batch\_size=1 and some other dimension=1
- `torch.squeeze()` removes all dimensions of 1

good practice: explicitly specify squeeze dimensions (**dim** argument)

```
[ ]  1  def forward(self, x):
    2      # x shape: 1 x 100 x 1
    3      y = x.squeeze()
    4
    5      # y shape: 100 (instead of 1 x 100)
```



# Wrong tensor shape - intermediate ops

```
[27] 1 class SimpleNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.conv_layer = nn.Conv2d(
5             in_channels=3,
6             out_channels=3,
7             kernel_size=1
8         )
9         self.lin = nn.Linear(in_features=3072, out_features=64)
10
11    def forward(self, x):
12        # x shape: batch_size x 3 x 32 x 32
13        # o1 shape: batch_size x 3 x 32 x 32
14        o1 = self.conv_layer(x)
15
16        # o2 shape: batch_size x 64
17        o2 = self.lin(o1)
18
19        return o2
```

```
[29] 1 net = SimpleNet()
2 batch_size = 10
3 x = torch.rand(batch_size, 3, 32, 32)
4 out = net(x)
5 print(out.shape)
```

scenario:

- *linear layer after conv2d layer*
- linear expects input=3x32x32=3072
- linear layer 

```
115    def forward(self, input: Tensor) -> Tensor:
--> 116        return F.linear(input, self.weight, self.bias)
117
118    def extra_repr(self) -> str:
```

RuntimeError: mat1 and mat2 shapes cannot be multiplied (960x32 and 3072x64)



# Wrong tensor shape - intermediate ops

```
▶ 1 class SimpleNet(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.conv_layer = nn.Conv2d(
5             in_channels=3,
6             out_channels=3,
7             kernel_size=1
8         )
9         self.lin = nn.Linear(in_features=3072, out_features=64)
10
11    def forward(self, x):
12        # x shape: batch_size x 3 x 32 x 32
13        # o1 shape: batch_size x 3 x 32 x 32
14        batch_size = x.size(0)
15        o1 = self.conv_layer(x)
16
17        # o1 shape: batch_size x 3072
18        o1 = o1.reshape(batch_size, -1)
19
20        # o2 shape: batch_size x 64
21        o2 = self.lin(o1)
22
23        return o2
```

```
[33] 1 net = SimpleNet()
2 batch_size = 10
3 x = torch.rand(batch_size, 3, 32, 32)
4 out = net(x)
5 print(out.shape)

torch.Size([10, 64])
```

scenario:

- *linear layer after conv2d layer*
- linear expects input=3x32x32=3072
- reshape input 



# Model summary

- tools such as [torchinfo](#) show expected input/output shapes for each layer

```
from torchinfo import summary

model = ConvNet()
batch_size = 16
summary(model, input_size=(batch_size, 1, 28, 28))
```

Layer (type:depth-idx)	Input Shape	Output Shape	Param #
SingleInputNet	[7, 1, 28, 28]	[7, 10]	--
-Conv2d: 1-1	[7, 1, 28, 28]	[7, 10, 24, 24]	260
-Conv2d: 1-2	[7, 10, 12, 12]	[7, 20, 8, 8]	5,020
-Dropout2d: 1-3	[7, 20, 8, 8]	[7, 20, 8, 8]	--
-Linear: 1-4	[7, 320]	[7, 50]	16,050
-Linear: 1-5	[7, 50]	[7, 10]	510

Total params: 21,840

Trainable params: 21,840

Non-trainable params: 0

Total mult-adds (M): 3.41

# Wrong tensor shape - good practices

- write down tensor shapes after each operation

```
# batch_size x 50
h1 = torch.rand(BATCH_SIZE, 50)
h2 = torch.rand(BATCH_SIZE, 50)

# batch_size x 100
h_cat = torch.cat((h1, h2), dim=1)
```

- use tools such as **torchinfo** to detect *potential layer mismatches*
- when debugging, *log all the tensor shapes* leading up to the shape error
- carefully read Pytorch documentation for layers (linear, conv2d etc), operations (concat, reshape, etc.) and loss functions

# Wrong tensor device

- key idea: most operations require Tensors **to be on the same device**
- this is also true for distributed training (multi-GPU training)
  - the distributed framework takes care of syncing gradients and parameters
  - local computation still has to be performed on the same GPU
- mixing devices (**CPU->GPU**, or different **GPUs**) results in runtime errors

```
-> 2846     return torch._C._nn.cross_entropy_loss(input, target, weight, _Reduction.get_enum(reduction), ignore_index,
label_smoothing)
2847
2848
```

**RuntimeError:** Expected all tensors to be on the same device, but found at least two devices, cuda:0 and cpu! (when checking argument for argument target in method wrapper\_nll\_loss\_forward)

# Wrong tensor device

common scenarios:

- model was moved to GPU but input data is still on CPU due to `.to()` operation

```
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')

# Module.to() is an in-place operation
model.to(device)

# Tensor.to() is not an in-place operation
images=images.to(device)]
```

- model was loaded from a checkpoint but we forgot to move it to GPU
- forgot to move labels to GPU

# Wrong tensor device - best practices

- write device agnostic code to easily switch between the two

```
device = torch.device('cuda') if torch.cuda.is_available() else torch.device('cpu')
```

- make sure model, inputs and labels **are on the same device**
  - remember DataLoader usually loads data on CPU
  - pay attention to `.to()` operations
  - log model/tensor device info

# Wrong tensor type

- PyTorch ops expect certain Tensor types (**float32** for operations, **long** for indices/labels)
- Mismatched types usually lead to runtime errors

```
/usr/local/lib/python3.7/dist-packages/torch/nn/functional.py in
 1846     if has_torch_function_variadic(input, weight, bias):
 1847         return handle_torch_function(linear, (input, weig
-> 1848     return torch._C._nn.linear(input, weight, bias)
 1849
 1850
```

RuntimeError: expected scalar type Float but found Double

# Wrong tensor type

input data and network mismatch:

- parameters in **torch.nn** layers have type **float32** (Float) by default
- NumPy arrays are converted to **float64** (Double) Tensors => incompatible with layers

```
1 x = torch.rand(3, 18)
2 print(x.dtype)
```

```
torch.float32
```

```
1 y = torch.tensor(np.array([1.07, 3.87, 4.32]))
2 print(y.dtype)
```

```
torch.float64
```

# Wrong tensor type

input data and network mismatch:

- pretrained model parameters have half-precision (**float16**)
- input is **float32**

```
1 model = MLP(input_size=3072, hidden_size=1024, activation_fn=nn.ReLU())
2 model.half() # convert parameters to float16
3 x = torch.rand(3072) # float32 input
4 out = model(x)
```

◆ 3 frames

[/usr/local/lib/python3.9/dist-packages/torch/nn/modules/linear.py](#) in forward(self, input)

```
112
113     def forward(self, input: Tensor) -> Tensor:
--> 114         return F.linear(input, self.weight, self.bias)
115
116     def extra_repr(self) -> str:
```

RuntimeError: mat1 and mat2 must have the same dtype

# Wrong tensor type

labels and loss mismatch:

- labels are read as ***Float*** Tensors
- we forget to convert them to ***Long*** Tensors, required by **`nn.CrossEntropyLoss`**

---

◆ 5 frames

```
/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py in nll_loss(input, target, weight, size_average, ignore_index,
2113                               .format(input.size(0), target.size(0)))
2114     if dim == 2:
-> 2115         ret = torch._C._nn.nll_loss(input, target, weight, _Reduction.get_enum(reduction), ignore_index)
2116     elif dim == 4:
2117         ret = torch._C._nn.nll_loss2d(input, target, weight, _Reduction.get_enum(reduction), ignore_index)
```

`RuntimeError`: expected scalar type `Long` but found `Double`

# Wrong tensor type - best practices

- inspect Tensor types:
  - use `tensor.dtype` to check types before key operations
- check documentation:
  - Loss Functions:
    - `nn.CrossEntropyLoss` expects **Float** outputs and **Long** labels
  - Indexing and Embedding layers:
    - `nn.Embedding` expects **Long** inputs (indices of words)
    - `torch.gather()` expects **Long** indices

# Wrong tensor range

- labels are stored as  $\{1, 2, \dots, C\}$  instead of  $\{0, 1, \dots, C-1\}$
- this throws an error in the cost function when accessing the score or label with index  $C$  (out of range)

Original Data

Team	Points
1	25
1	12
2	15
2	14
2	19
2	23
3	25
3	29

Label Encoded Data

Team	Points
0	25
0	12
1	15
1	14
1	19
1	23
2	25
2	29



# Wrong tensor range

- when working with text data, each word has a unique index  $\{0..V-1\}$
- for  $V = 10000$ ,  $x[10000]$  results in an ‘index out of range’ error

## Vocabulary



[source](#)

B

# Out of memory errors



Suhail ✅

@Suhail

...

This is currently the bane of my existence:

"RuntimeError: CUDA out of memory. Tried to allocate 220.00 MiB (GPU 0; 15.78 GiB total capacity; 13.16 GiB already allocated; 201.44 MiB free; 14.44 GiB reserved in total by PyTorch)"

This must be an AI rite of passage.

3:41 AM · Jul 20, 2022

B

# Out of memory errors - causes

- memory inefficiency:
  - tensors are not detached (either outputs or losses)

# Out of memory errors - causes

- memory inefficiency:
  - tensors are not detached (either outputs or losses)
- large batch size
  - increases memory requirements

# Out of memory errors - causes

- memory inefficiency:
  - tensors are not detached (either outputs or losses)
- large batch size
  - increases memory requirements
- large model:
  - model size:
    - 7B model needs 14GB VRAM memory just to store its FP16 weights
  - training overhead:
    - need to store activations/gradients/optimizer states => a 7B model requires ~90-100GB VRAM memory

# Out of memory errors - test time

- check for memory inefficiencies:
  - wrap forward code in `torch.no_grad()`, which disables gradient calculation and reduces memory consumption
  - if you save outputs (segmentation maps, probabilities) for later inspection, move Tensors to CPU

# Out of memory errors - test time

- check for memory inefficiencies:
  - wrap forward code in `torch.no_grad()`, which disables gradient calculation and reduces memory consumption
  - if you save outputs (segmentation maps, probabilities) for later inspection, move them to CPU
- reduce precision:
  - convert model to FP16/BF16 if supported
  - quantize model (INT8)

# Out of memory errors - test time

- check for memory inefficiencies:
  - wrap forward code in `torch.no_grad()`, which disables gradient calculation and reduces memory consumption
  - if you save outputs (segmentation maps, probabilities) for later inspection, move them to CPU
- reduce precision:
  - convert model to FP16/BF16 if supported
  - quantize model (INT8)
- reduce input:
  - reduce batch size
  - feed shorter sequences (if working with text/time series etc.)

# **Out of memory errors - train time**

- check for memory leaks (detach loss from computational graph before storing it)

# Out of memory errors - train time

- check for memory leaks (detach loss from computational graph before storing it)
- reduce model training footprint:
  - train with mixed precision (FP16 + FP32)
  - train with parameter-efficient methods (LoRA)
  - reduce model size (less layers etc.) if training from scratch

# Out of memory errors - train time

- check for memory leaks (detach loss from computational graph before storing it)
- reduce model training footprint:
  - train with mixed precision (FP16 + FP32)
  - train with parameter-efficient methods (LoRA)
  - reduce model size (less layers etc.) if training from scratch
- reduce batch size:
  - use smaller BS with gradient accumulation to simulate larger batches

# Out of memory errors - train time

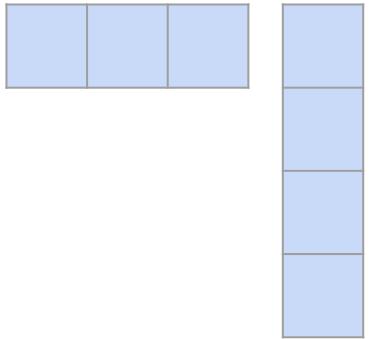
- check for memory leaks (detach loss from computational graph before storing it)
- reduce model training footprint:
  - train with mixed precision (FP16 + FP32)
  - train with parameter-efficient methods (LoRA)
  - reduce model size (less layers etc.) if training from scratch
- reduce batch size:
  - use smaller BS with gradient accumulation to simulate larger batches
- activation checkpointing:
  - recompute activations during backprop to save memory

# Out of memory errors - train time

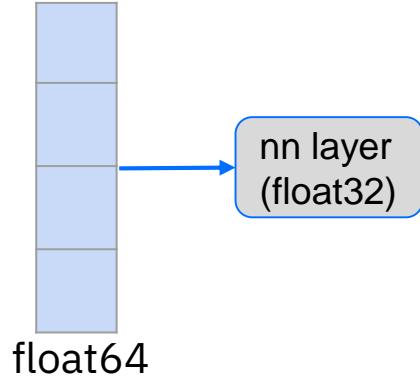
- check for memory leaks (detach loss from computational graph before storing it)
- reduce model training footprint:
  - train with mixed precision (FP16 + FP32)
  - train with parameter-efficient methods (LoRA)
  - reduce model size (less layers etc.) if training from scratch
- reduce batch size:
  - use smaller BS with gradient accumulation to simulate larger batches
- activation checkpointing:
  - recompute activations during backprop to save memory
- if more GPUs available, model parallelism:
  - split model to several GPUs

# Syntactic Issues - recap

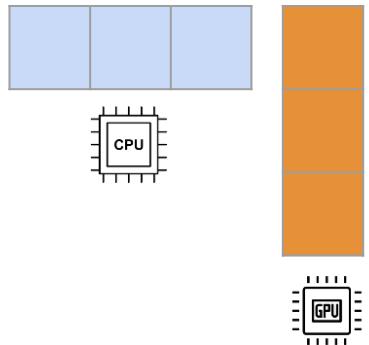
shape errors



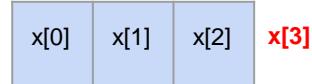
type errors



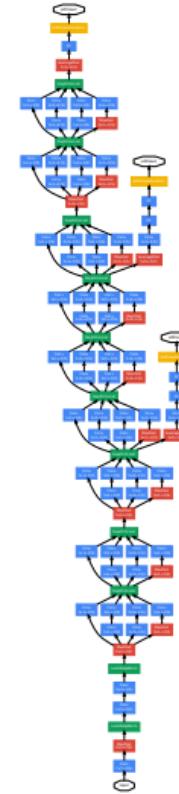
device errors



range errors



OOM errors



# Machine Learning bugs

syntactic

```
-----  
ValueError                                Traceback (most recent call last)  
<ipython-input-4-59c95019453c> in <module>()  
      1 model = Sequential()  
----> 2 model.add(Conv1D(filters=64, kernel_size=3, activation='relu', input_shape=(50,)))  
      3 model.add(Conv1D(filters=64, kernel_size=3, activation='relu'))  
      4 model.add(Dropout(0.5))  
      5 model.add(MaxPooling1D(pool_size=2))  
  
~/anaconda3/lib/python3.6/site-packages/keras/engine/sequential.py in add(self, layer)  
    163         # and create the node connecting the current layer  
    164         # to the input layer we just created.  
--> 165             layer(x)  
    166             set_inputs = True  
    167         else:  
  
~/anaconda3/lib/python3.6/site-packages/keras/engine/base_layer.py in __call__(self, inputs, **kwargs)  
    412     # Raise exceptions in case the input is not compatible  
    413     # with the input_spec specified in the layer constructor.  
--> 414     self.assert_input_compatibility(inputs)  
    415  
    416     # Collect input shapes to build layer.  
  
~/anaconda3/lib/python3.6/site-packages/keras/engine/base_layer.py in assert_input_compatibility(self, inputs)  
    309             self.name + ': expected ndim=' +  
    310                 str(spec.ndim) + ', found ndim=' +  
--> 311                 str(K.ndim(x)))  
    312         if spec.max_ndim is not None:  
    313             ndim = K.ndim(x)  
  
ValueError: Input 0 is incompatible with layer conv1d_1: expected ndim=3, found ndim=2
```



functional



# Machine Learning functional bugs

- no syntax errors present
- however, model either:
  - fails to learn the train data
  - doesn't generalize to new data



# Functional bugs - sources of error

- **logic errors** (incorrect arguments, wrong labels)
- **dataset issues** (data preprocessing, augmentation, train/test mismatch)
- **hyperparameter misconfigurations** (learning rates, regularization issues)



# Logic errors

normalization/standardization is:

- missing
- computed over the wrong dimension

Age	Salary	Volume
28	26,000	4
22	32,000	10
25	29,000	86
24	38,000	2
26	27,000	71
22	42,000	88
25	41,000	90
23	26,500	72
22	29,500	14
24	36,000	8

# Logic errors

normalization/standardization is:

- missing
- computed over the wrong dimension

Age	Salary	Volume
28	26,000	4
22	32,000	10
25	29,000	86
24	38,000	2
26	27,000	71
22	42,000	88
25	41,000	90
23	26,500	72
22	29,500	14
24	36,000	8

[ [0.001 1. 0. ]  
[0. 1. 0. ]  
[0. 1. 0.002]  
[0.001 1. 0. ]  
[0. 1. 0.002]  
[0. 1. 0.002]  
[0. 1. 0.002]  
[0. 1. 0.002]  
[0. 1. 0. ]  
[0. 1. 0. ] ]



# Logic errors

normalization/standardization is:

- missing
- computed over the wrong dimension

Age	Salary	Volume
28	26,000	4
22	32,000	10
25	29,000	86
24	38,000	2
26	27,000	71
22	42,000	88
25	41,000	90
23	26,500	72
22	29,500	14
24	36,000	8



```
[[1.    , 0.    , 0.023],  
 [0.    , 0.375, 0.091],  
 [0.5   , 0.188, 0.955],  
 [0.333, 0.75  , 0.    ],  
 [0.667, 0.062, 0.784],  
 [0.    , 1.    , 0.977],  
 [0.5   , 0.938, 1.    ],  
 [0.167, 0.031, 0.795],  
 [0.    , 0.219, 0.136],  
 [0.333, 0.625, 0.068]])
```



# Logic errors

softmax is:

applied over the wrong dimension

```
BATCH_SIZE = 8
# batch_size x 2
output_scores = torch.rand(BATCH_SIZE, 2)
softmax = nn.Softmax(dim=0)
# batch_size x 2
output_probabilities = softmax(output_scores)
print(output_probabilities)
```

```
tensor([[0.1392, 0.1821],
       [0.1809, 0.0744],
       [0.1111, 0.1447],
       [0.1298, 0.0712],
       [0.0753, 0.0696],
       [0.0687, 0.1845],
       [0.1725, 0.1599],
       [0.1225, 0.1135]])
```

# Logic errors

softmax is:

applied over the wrong dimension

```
BATCH_SIZE = 8  
# batch_size x 2  
output_scores = torch.rand(BATCH_SIZE, 2)  
softmax = nn.Softmax(dim=0)  
# batch_size x 2  
output_probabilities = softmax(output_scores)  
print(output_probabilities)
```

```
tensor([[0.1392, 0.1821],  
       [0.1809, 0.0744],  
       [0.1111, 0.1447],  
       [0.1298, 0.0712],  
       [0.0753, 0.0696],  
       [0.0687, 0.1845],  
       [0.1725, 0.1599],  
       [0.1225, 0.1135]])
```

unnecessary

```
loss_fn = nn.CrossEntropyLoss()  
y = torch.tensor([1, 0, 0, 1, 1, 1, 0, 0])  
  
# batch_size x 2  
output_probabilities = softmax(output_scores)  
  
loss = loss_fn(output_probabilities, y)
```

*nn.CrossEntropyLoss()* requires  
raw unnormalized scores as input



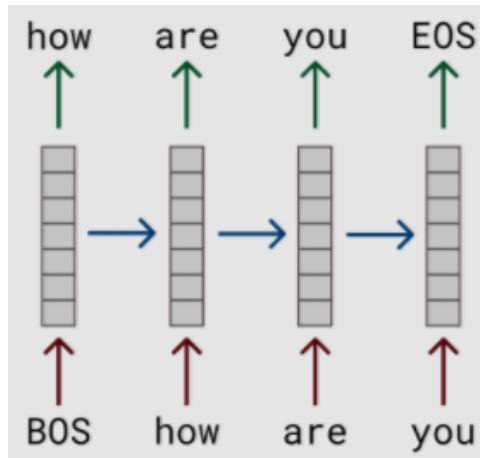
# Logic errors

train/evaluation mode mismatch:

- when model is set in *eval()* mode, some operations behave differently (*Dropout, Batch Normalization* etc.)
- **common mistake:** after evaluation, model is **NOT** set back to *train mode*, which results in unwanted behaviour:
  - dropout **remains disabled** during training => no more regularization
  - batch normalization uses **wrong estimates of mean and variance** => bad updates

# Logic errors

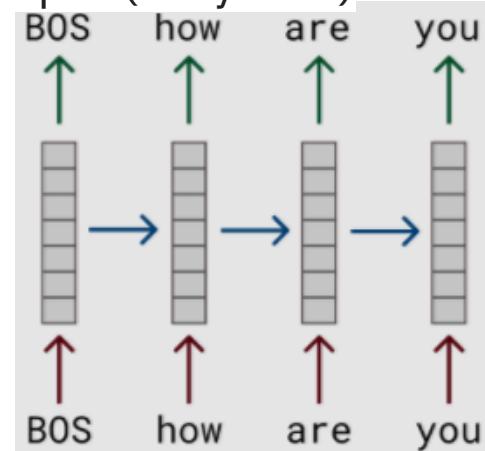
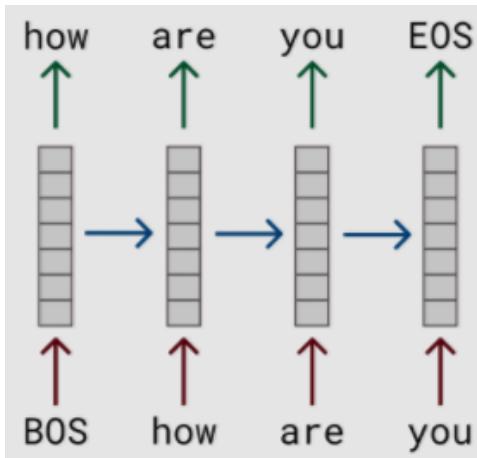
- in language models, for each input word you have to predict the *next* word
- labels = inputs shifted by one to the left



BOS = beginning of sentence, EOS = end of sentence

# Logic errors

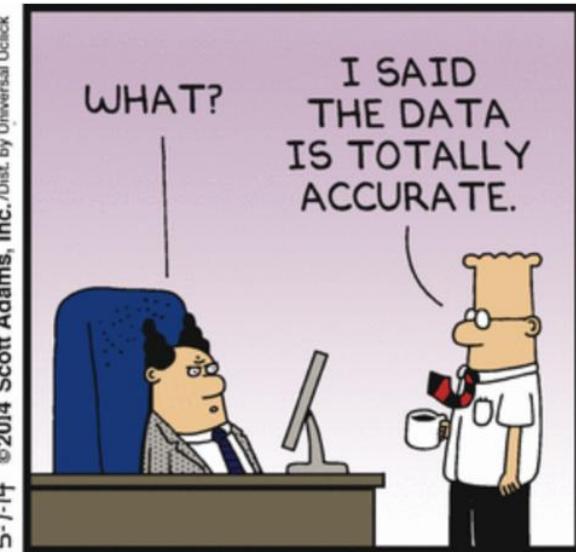
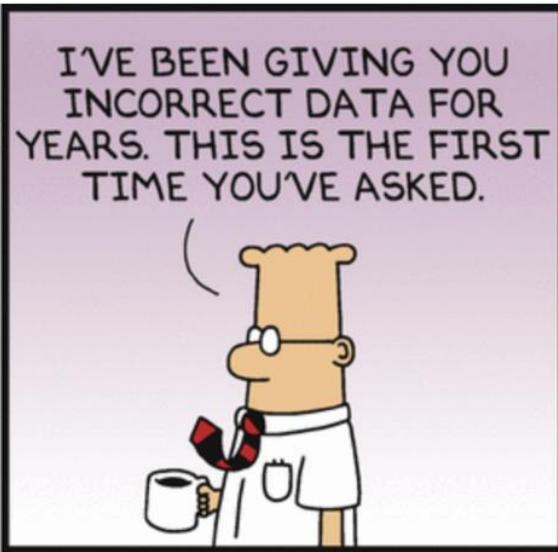
- in language models, for each input word you have to predict the *next* word
- labels = inputs shifted by one to the left
- if labels are not shifted properly, model solves different task by mistake
- in this example, model predicts its own input (easy task)



BOS = beginning of sentence, EOS = end of sentence

# Dataset issues

check if data is faulty: mislabeled, unshuffled, improperly augmented etc.



# Dataset issues

- **NaN** (not a number) value appears in input data by mistake (missing/corrupted data, preprocessing error etc.)
- **no syntax errors** raised during:
  - feedforward
  - loss computation
  - backpropagation

```
BATCH_SIZE = 8
model = SimpleMLP()
loss_fn = nn.CrossEntropyLoss()
y = torch.tensor([1, 0, 0, 1, 1, 0, 0, 0])

# batch_size x 10
input = torch.rand(BATCH_SIZE, 10)
input[0,0] = float('nan')

# batch_size x 2
output_scores = model(input)
loss = loss_fn(output_scores, y)
print(loss)

tensor(nan, grad_fn=<NllLossBackward>)

loss.backward()
for p in model.parameters():
    print(p.grad)

tensor([[nan, nan, nan, nan, nan, nan, nan, nan, nan, nan],
```

# Dataset issues

- use `torch.isnan(x).any()` to manually check tensors
- use `torch.autograd.detect_anomaly()`

```
▶ 1  with autograd.detect_anomaly():
  2      x = torch.tensor([[1.0, 2.0], [3.0, float('nan')]], requires_grad=True)
  3      y = torch.tensor([[1.0, 2.0], [3.0, 4.0]], requires_grad=True)
  4      z = x * y
  5      out = torch.sum(z)
  6      out.backward()

  4      z = x * y
  5      out = torch.sum(z)
----> 6      out.backward()

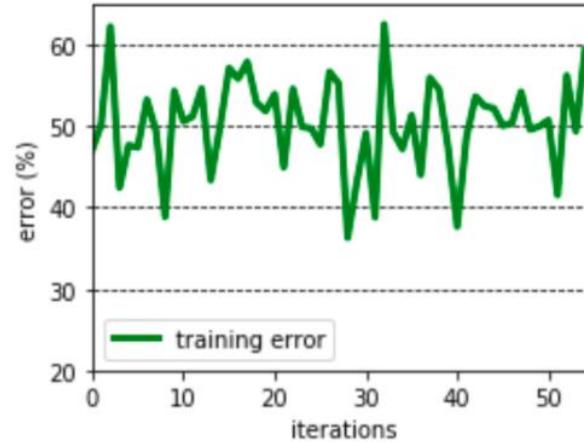
----- ◇ 1 frames -----
/usr/local/lib/python3.10/dist-packages/torch/autograd/ __init__.py in backward(tensors, grad_tensors, retai
  264     # some Python versions print out the first line of a multi-line function
  265     # calls in the traceback and some print out the last line
--> 266     Variable._execution_engine.run_backward( # Calls into the C++ engine to run the backward pass
  267         tensors,
  268         grad_tensors_,

RuntimeError: Function 'MulBackward0' returned nan values in its 1th output.
```

# Dataset issues

- examples and labels are read from individual files
- what is the issue?

```
1 features = glob.glob('path/to/features/*')
2 labels = glob.glob('path/to/labels/*')
3 train(features, labels)
```

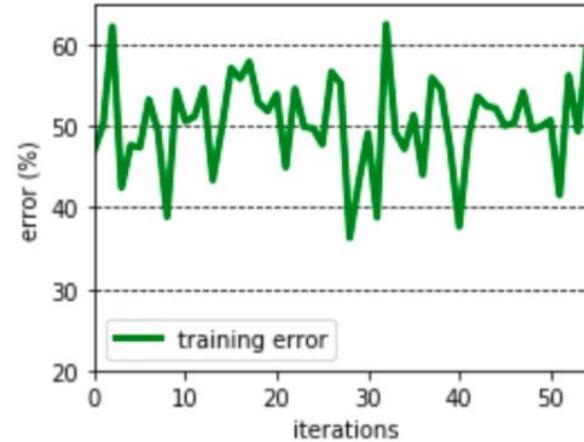


source: <https://fullstackdeeplearning.com/spring2021/lecture-7/>

# Dataset issues

- examples and labels are read from individual files
- what is the issue?
  - files are read in random order, `features[i]` and `labels[i]` are not related

```
1 features = glob.glob('path/to/features/*')
2 labels = glob.glob('path/to/labels/*')
3 train(features, labels)
```



source: <https://fullstackdeeplearning.com/spring2021/lecture-7/>

# Dataset issues

what is the issue?

```
# instantiate dataloader
train_dataloader = DataLoader(
    train_dataset,
    batch_size=BATCH_SIZE
)
```

# Dataset issues

what is the issue?

- training examples are not shuffled

```
# instantiate dataloader
train_dataloader = DataLoader(
    train_dataset,
    batch_size=BATCH_SIZE
)
```



```
# instantiate dataloader
train_dataloader = DataLoader(
    train_dataset,
    batch_size=BATCH_SIZE,
    shuffle=True,
)
```



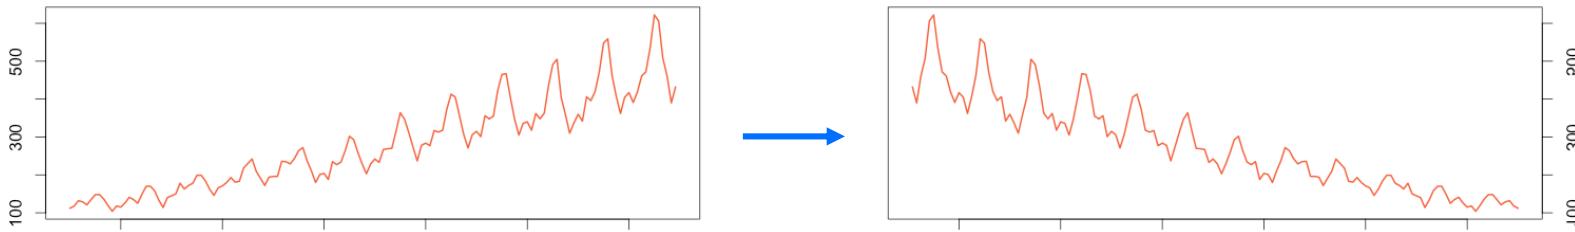
B

# Dataset issues

- dataset augmentation commonly improves performance
- however, some augmentations are not meaningful and actually hurt model's performance:
  - flipping/rotating logos may not improve logo recognition



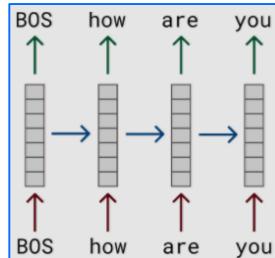
- flipping time series changes trends



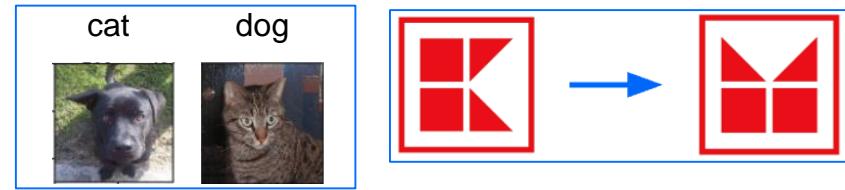
# Functional Issues - recap

1. fix **logical errors** (*bad normalization, softmax over the wrong dimension, wrong labels fed to the model etc.*)

Age	Salary	Volume
28	26,000	4
22	32,000	10
25	29,000	86
24	38,000	2
26	27,000	71
22	42,000	88



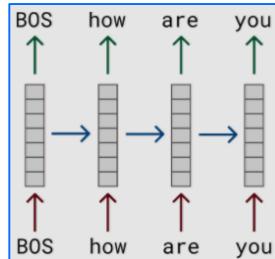
2. fix **dataset issues** (*data not shuffled, mislabeled data, improperly augmented etc.*)



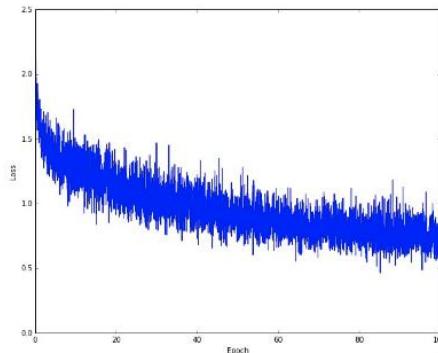
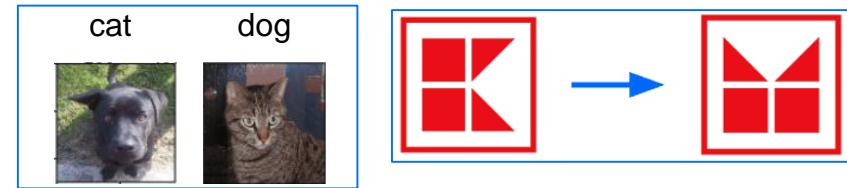
# Functional Issues - recap

1. fix **logical errors** (*bad normalization, softmax over the wrong dimension, wrong labels fed to the model etc.*)

Age	Salary	Volume
28	26,000	4
22	32,000	10
25	29,000	86
24	38,000	2
26	27,000	71
22	42,000	88



2. fix **dataset issues** (*data not shuffled, mislabeled data, improperly augmented etc.*)



3. we can now start looking at **hyperparameter misconfigurations**

# Debugging plan

1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)

# Debugging plan

1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)
2. overfit model on a few (1-2) examples
  - o if training loss ~0, start training on full dataset

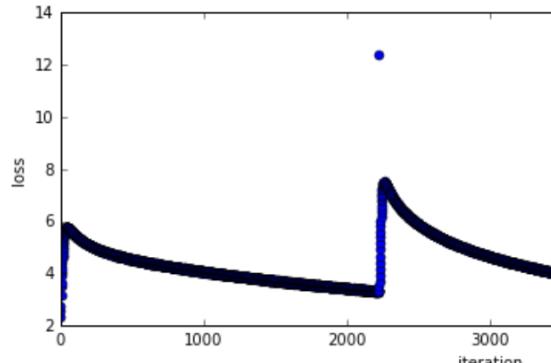
# Debugging plan

1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)
2. overfit model on a few (1-2) examples
  - o if training loss ~0, start training on full dataset
  - o if not, check for **logical and dataset errors**:
    - manually check input:
      - lots of zeros or NaN (*preprocessing error, corrupted data*)
      - duplicated rows (*dataset contains duplicates, augmentation bug*)
      - all batch examples have the same label (*no shuffling*)
    - check the arguments:
      - softmax taken over the correct **dimensions**
      - loss function receives the correct output (*unnormalized or normalized scores*)

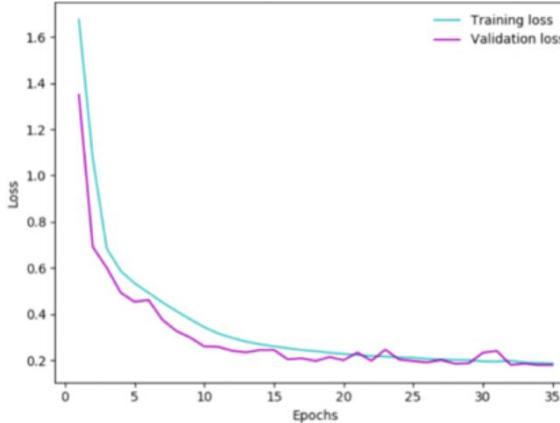
# Debugging plan

1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)
2. overfit model on a few (1-2) examples
  - if training loss ~0, start training on full dataset
  - if not, check for **logical and dataset errors**:
    - manually check input:
      - lots of zeros or NaN (*preprocessing error, corrupted data*)
      - duplicated rows (*dataset contains duplicates, augmentation bug*)
      - all batch examples have the same label (*no shuffling*)
    - check the arguments:
      - softmax taken over the correct **dimensions**
      - loss function receives the correct output (*unnormalized or normalized scores*)
3. start training on the full dataset
  - inspect **loss/metric learning curves** and look for **hyperparameters misconfigurations**
  - we may still have invisible bugs at this point! (*train data leaks into validation, bug when measuring validation error etc.*)

# Loss curves



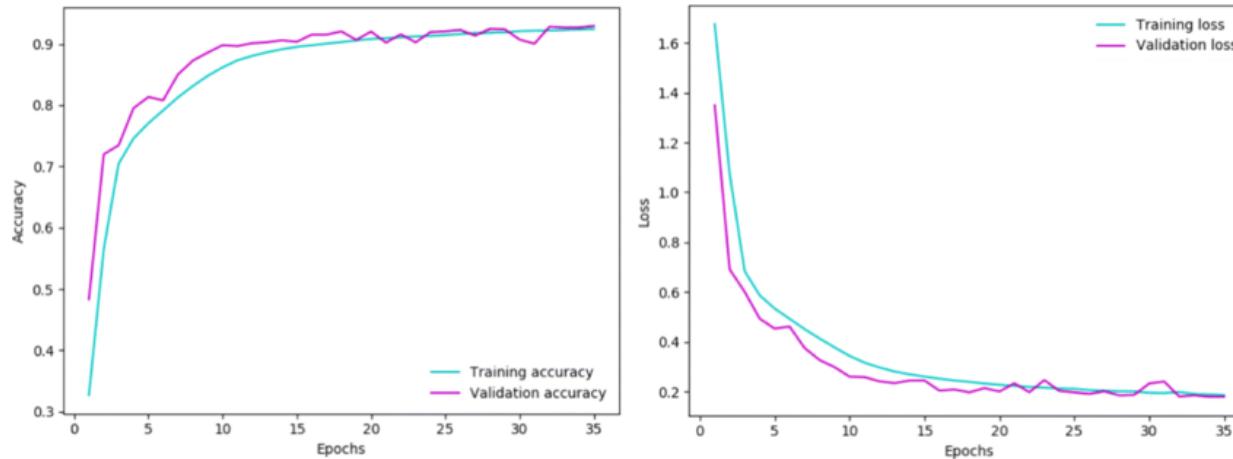
[ins.tumblr.com](https://ins.tumblr.com)



B

# Learning curves

- plot values for train/validation set over epochs/optimization steps:
  - **loss curves**: captures the optimization progress
  - **metric curves**: reflects our business/research goal (accuracy/F1/AUROC etc.)
- ideally, loss inversely correlated with metric



[source](#)

# Why look at learning curves?

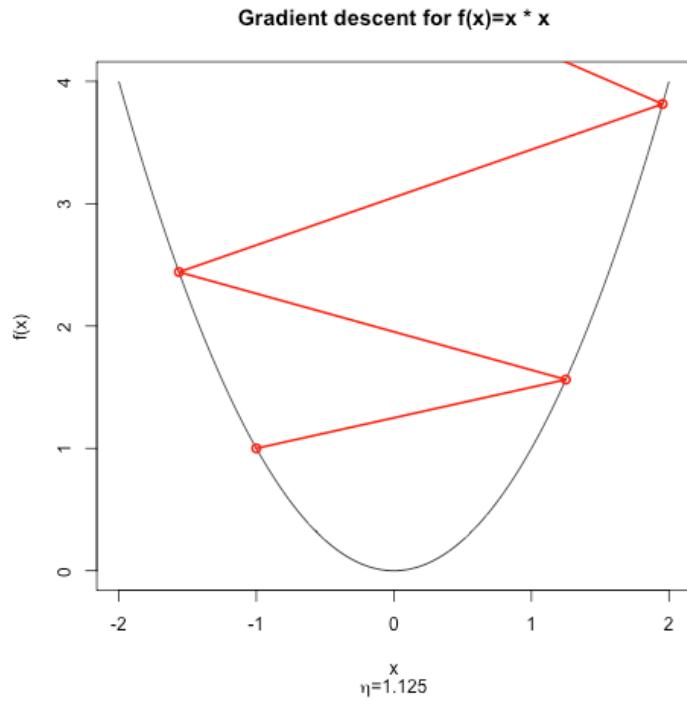
- detect if the model **learns the training data and/or generalizes to new data**
  - train and validation losses decrease at similar rate or show very small gap
  - if not, check for **hyperparameter misconfiguration**
- pick the best model based on the validation loss
- check if loss improvements **translate** to better metrics

# Does the model learn the train data?

- training loss increases
- training loss decreases slowly
- training loss decreases then stagnates
- training loss decreases

# Training loss increases

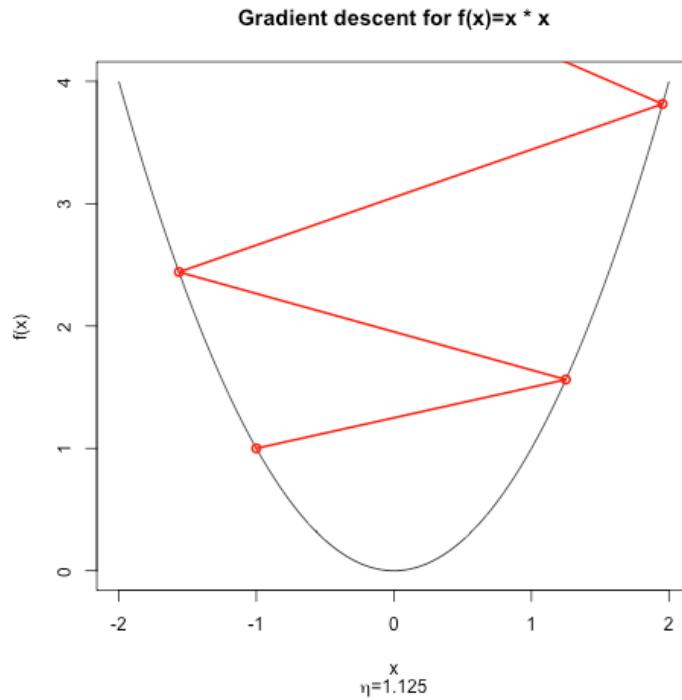
- Q: what does it usually indicate?



<https://stats.stackexchange.com/questions/364360/how-can-change-in-cost-function-be-positive>

# Training loss increases

- Q: what does it usually indicate?  
A: the learning rate may be *too high*, moving the parameters away from the local minimum

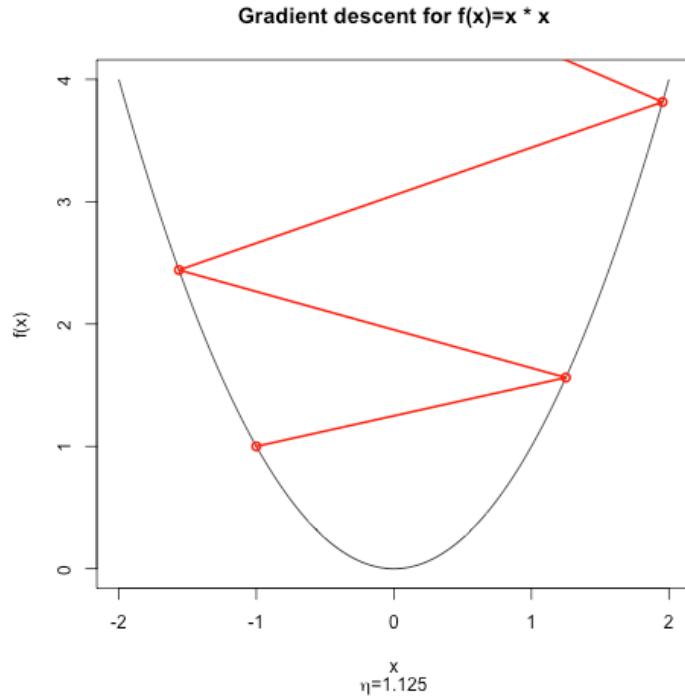


<https://stats.stackexchange.com/questions/364360/how-can-change-in-cost-function-be-positive>



# Training loss increases

- Q: what does it usually indicate?  
A: the learning rate may be *too high*, moving the parameters away from the local minimum
- **solution:** decrease learning rate



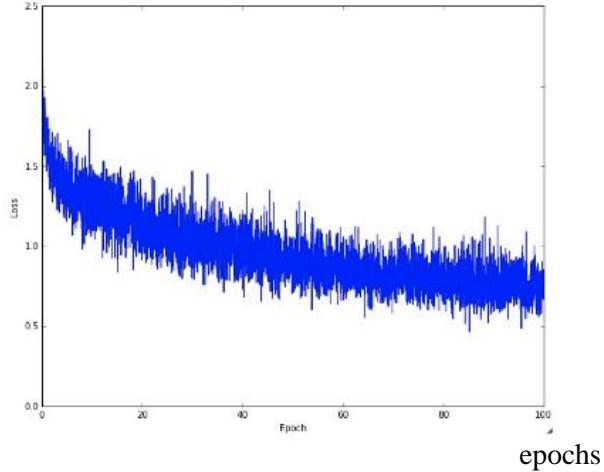
<https://stats.stackexchange.com/questions/364360/how-can-change-in-cost-function-be-positive>



# Training loss decreases slowly

- Q: what could it mean?

training loss

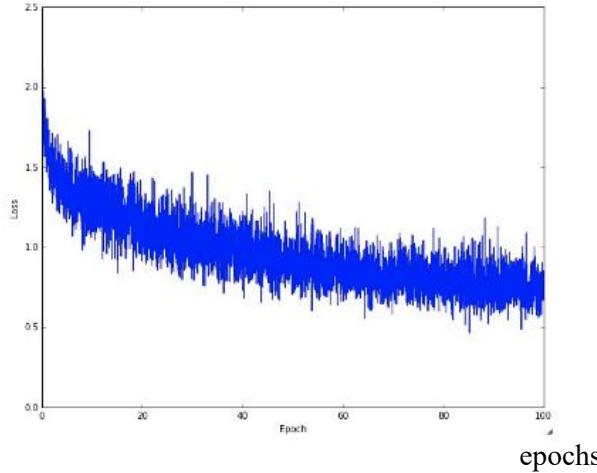


source: <https://cs231n.github.io/neural-networks-3/>

# Training loss decreases slowly

- Q: what could it mean?
- A: it could indicate a learning rate that is *too low* (loss decreases linearly)

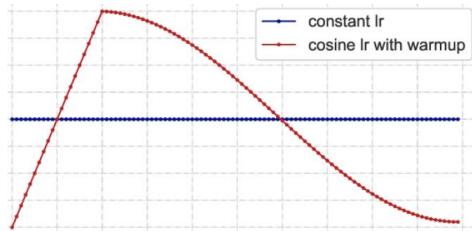
training loss



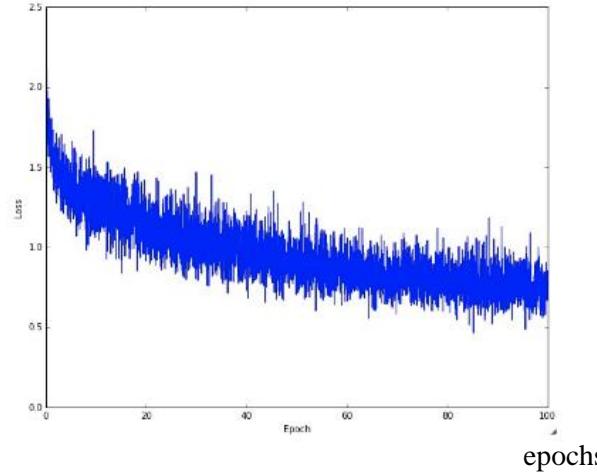
source: <https://cs231n.github.io/neural-networks-3/>

# Training loss decreases slowly

- Q: what could it mean?
- A: it could indicate a learning rate that is *too low* (loss decreases linearly)
- solution:
  - higher learning rate
  - use learning rate scheduler with warmup (LR increases during warmup then decreases)
    - [cosineLRWithWarmup](#): increase LR linearly then decrease with cosine



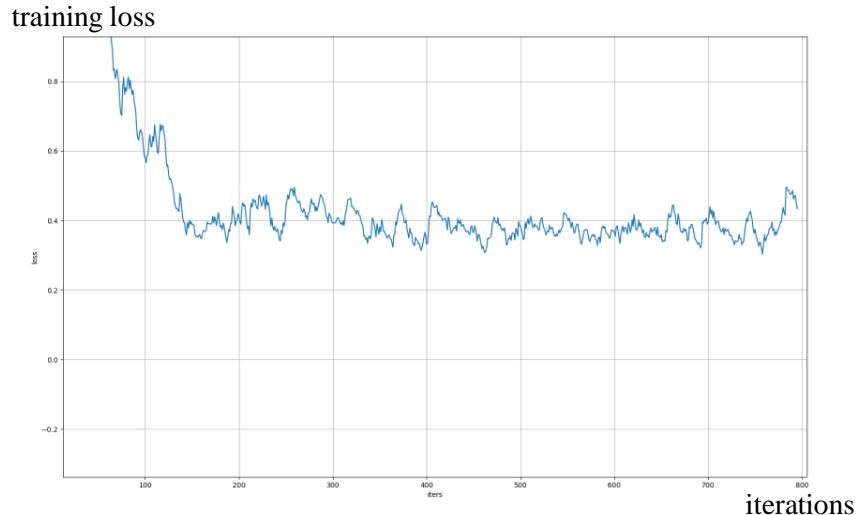
training loss



source: <https://cs231n.github.io/neural-networks-3/>

# Training loss decreases then stagnates

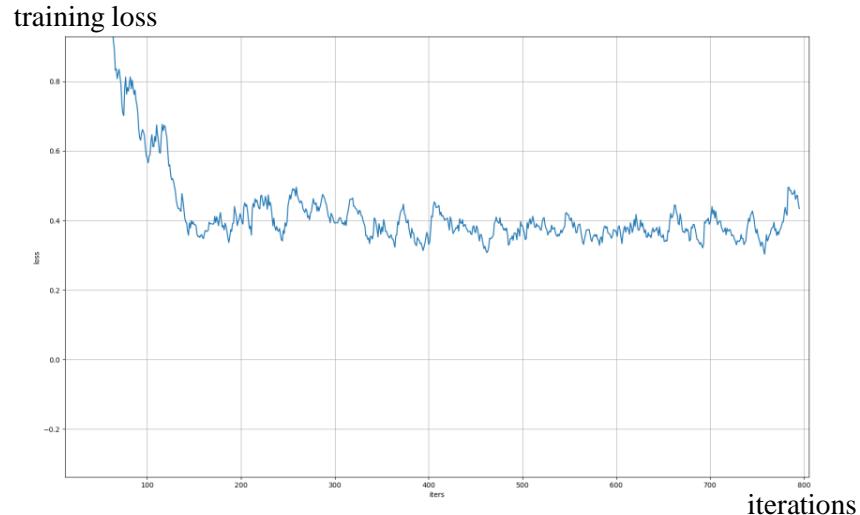
- Q: what could it mean?



source: <https://jumabek.wordpress.com/2017/03/21/781/>

# Training loss decreases then stagnates

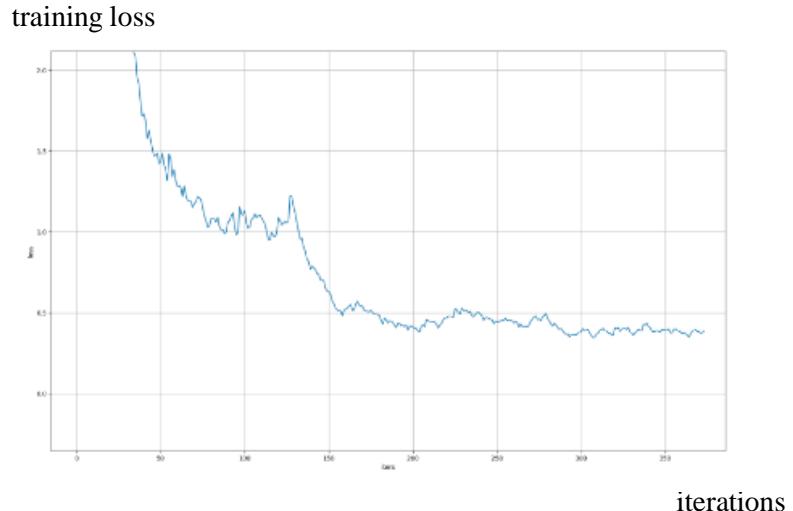
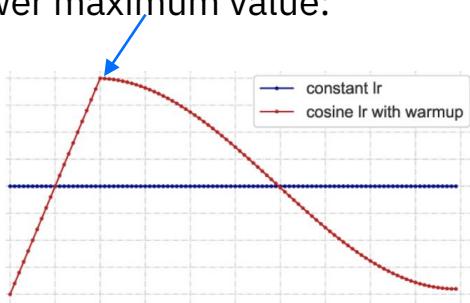
- Q: what could it mean?
- A: it may still indicate a learning rate that is too high (particularly if it oscillates a lot)



source: <https://jumabek.wordpress.com/2017/03/21/781/>

# Training loss decreases then stagnates

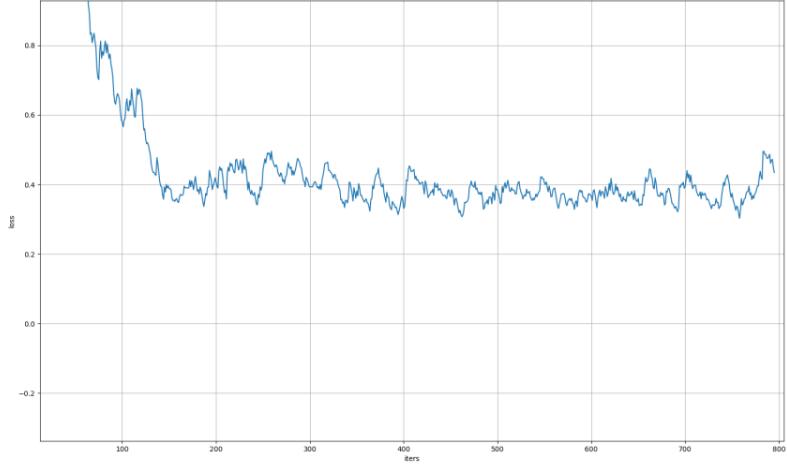
- Q: what could it mean?
- A: it may still indicate a learning rate that is too *high* (particularly if it oscillates a lot)
- solution:
  - start with a smaller learning rate ( $\frac{1}{2}$ - $\frac{1}{3}$  of original)
  - if using LR scheduler with warmup, try to lower maximum value:



source: <https://jumabek.wordpress.com/2017/03/21/781/>

# Training loss decreases then stagnates

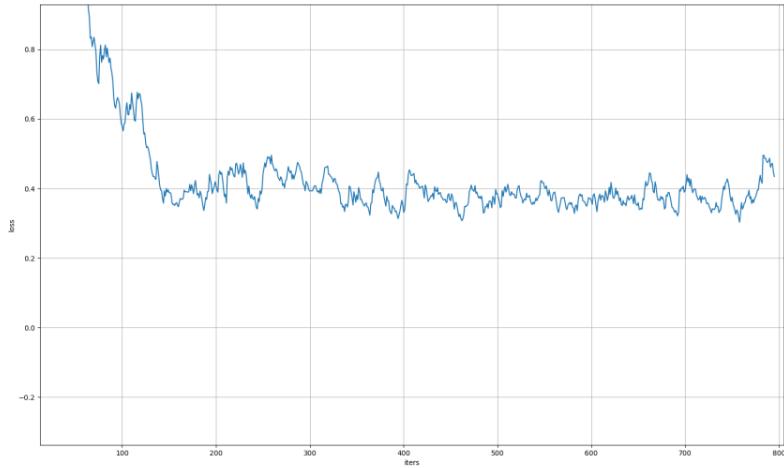
- what if higher LR still doesn't solve the issue?



source: <https://jumabek.wordpress.com/2017/03/21/781/>

# Training loss decreases then stagnates

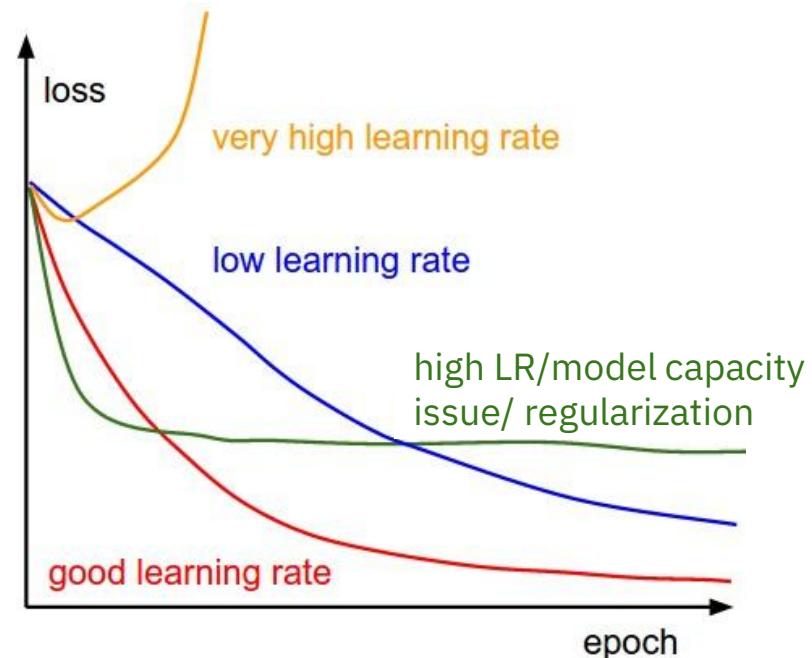
- what if higher LR still doesn't lower the train loss?
- possible issues:
  - the model is too small => *increase model capacity* (more layers, more neurons per layer)
  - it is heavily regularized => *reduce regularization* (smaller Dropout rate, smaller L2 coefficient)
  - lots of frozen parameters => unfreeze some of them



source: <https://jumabek.wordpress.com/2017/03/21/781/>

# training loss curves - recap

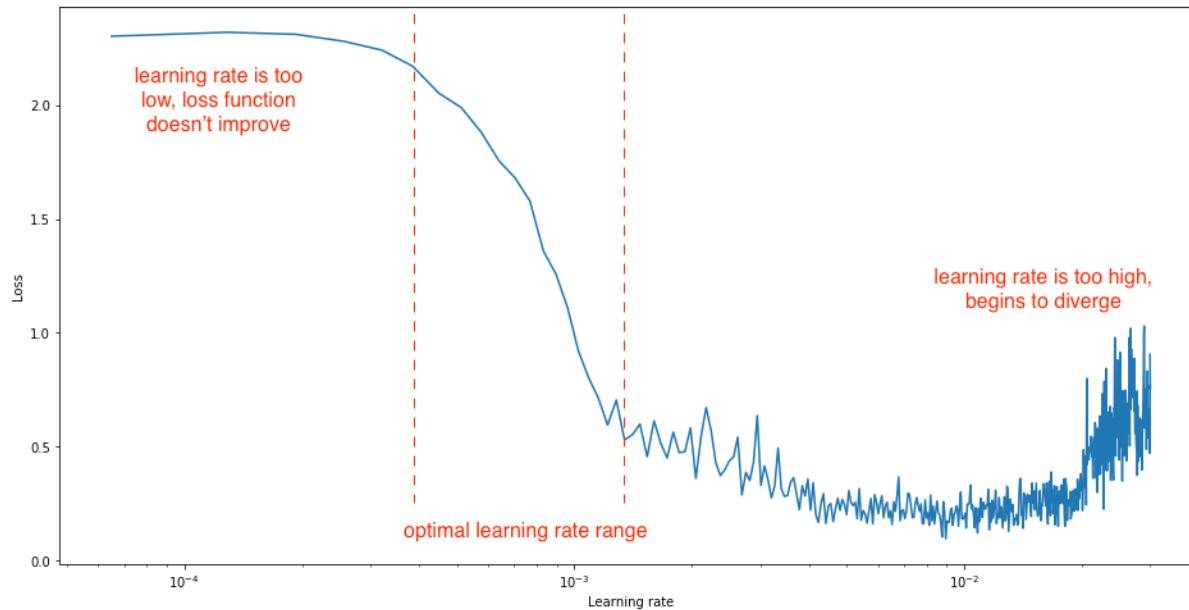
- loss increases early in the training (*high learning rate*)
- loss decreases slowly (potentially *low learning rate*)
- loss decreases quickly then reaches a plateau (*high learning rate or model capacity issues*)



source: <https://cs231n.github.io/neural-networks-3/>

# How to pick a good learning rate range?

- let model train and decrease learning rate for each new batch
- a good learning rate range:
  - starts when loss begins to drop steeply
  - ends when loss slows down, wiggles or starts increasing



source: <https://www.jeremyjordan.me/nn-learning-rate/>

# Learning curves scenarios

- training loss increases
- training loss decreases slowly
- training loss decreases then stagnates
- training loss decreases

we have looked at possible issues with training curves!

# Learning curves scenarios

- training loss increases
- training loss decreases slowly
- training loss decreases then stagnates
- training loss decreases, but:

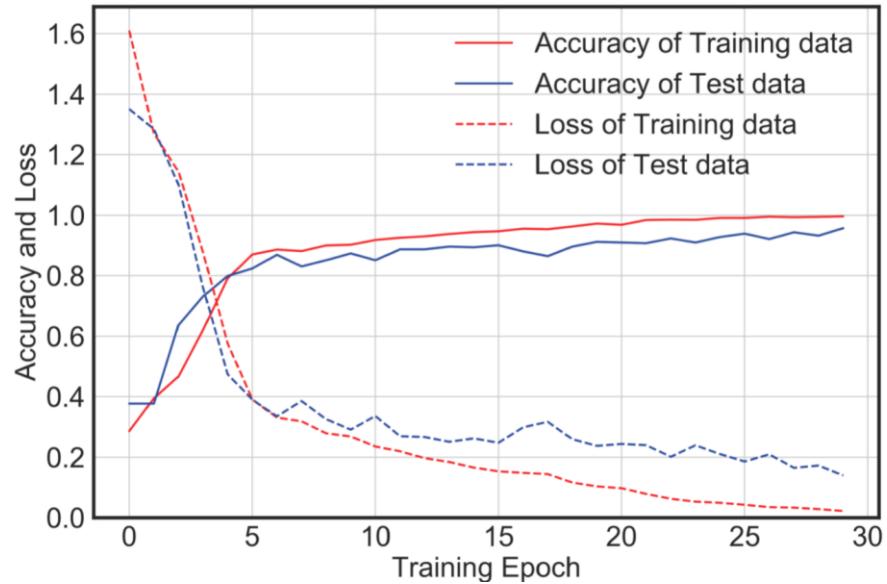
- validation loss  $\approx$  training loss
- validation loss  $>>$  training loss
- validation loss  $\approx$  training loss, then diverges
- validation loss  $<$  training loss
- validation loss  $<<$  training loss
- both validation loss and validation accuracy increase

however, we also want to see the generalization capability of the model  
=> inspect validation curves!



# Validation loss $\approx$ training loss

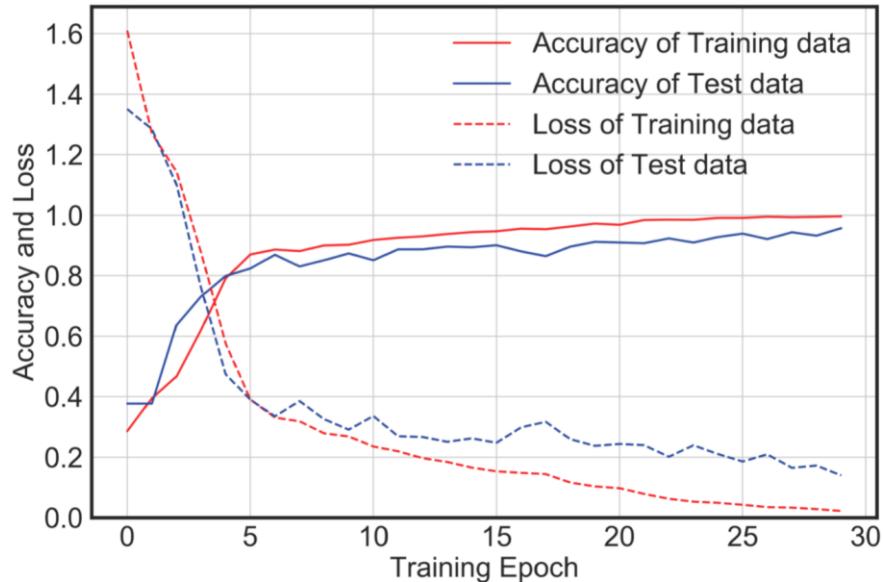
- Q: overfitting or underfitting?



source: <https://www.mdpi.com/1424-8220/20/9/2710/htm>

# Validation loss $\approx$ training loss

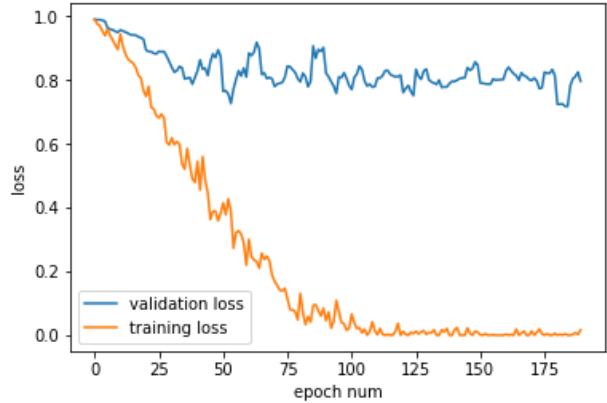
- Q: overfitting or underfitting?
- A: good fit
  - training and validation loss approach 0, with small gap between them
  - validation loss usually (but not always) slightly higher than training loss
  - training and validation accuracy also increase simultaneously (they are not always correlated with loss curves)



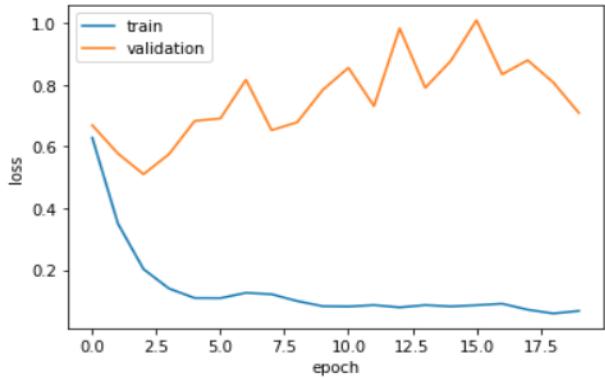
source: <https://www.mdpi.com/1424-8220/20/9/2710/htm>

# Validation loss >> training loss

- Q: overfitting or underfitting?



[source](#)

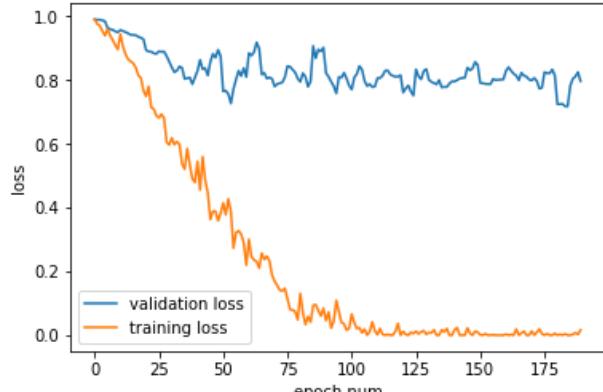


[source](#)

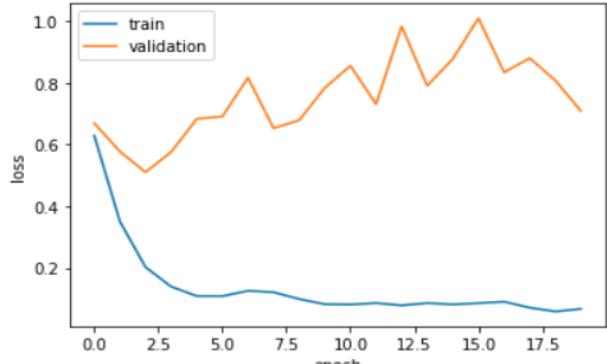


# Validation loss >> training loss

- Q: overfitting or underfitting?
- A: model is overfitting
  - training loss decreases, while validation loss either gets worse or doesn't improve
  - model learns patterns that are not useful outside training set
- solutions:
  - add regularization (Dropout/L2)
  - add more training data
  - reduce model size (fewer layers/neurons per layer)



[source](#)

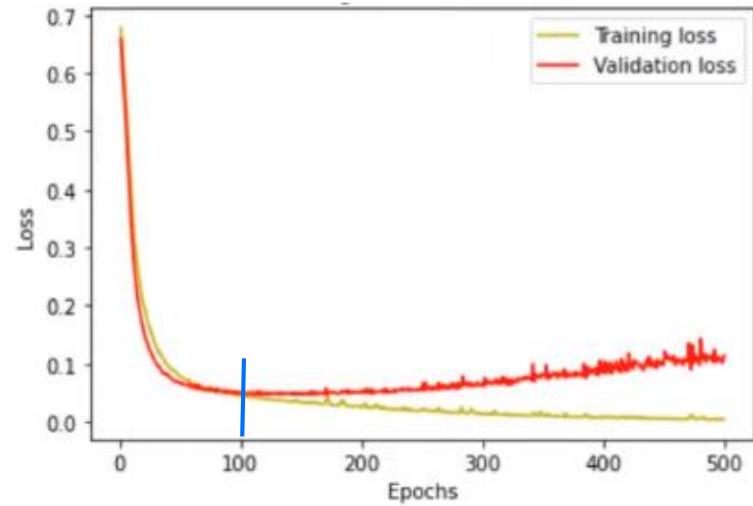


[source](#)



# Validation loss tracks train loss then diverges

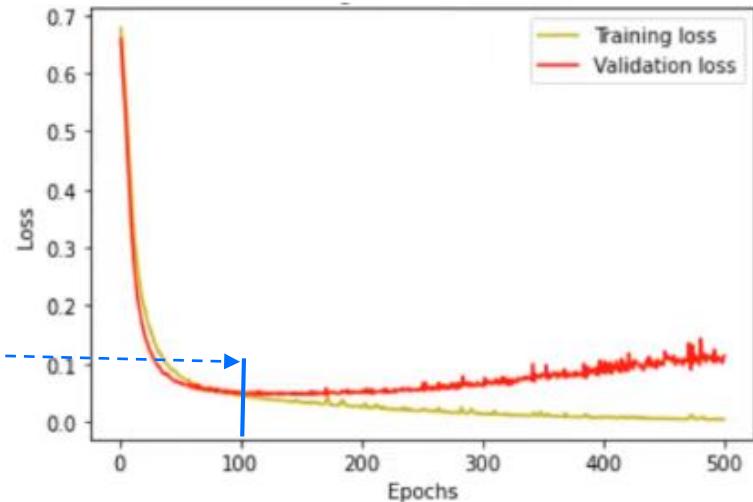
- Q: overfitting or underfitting?



source: <https://www.youtube.com/watch?v=p3CcfljycBA>

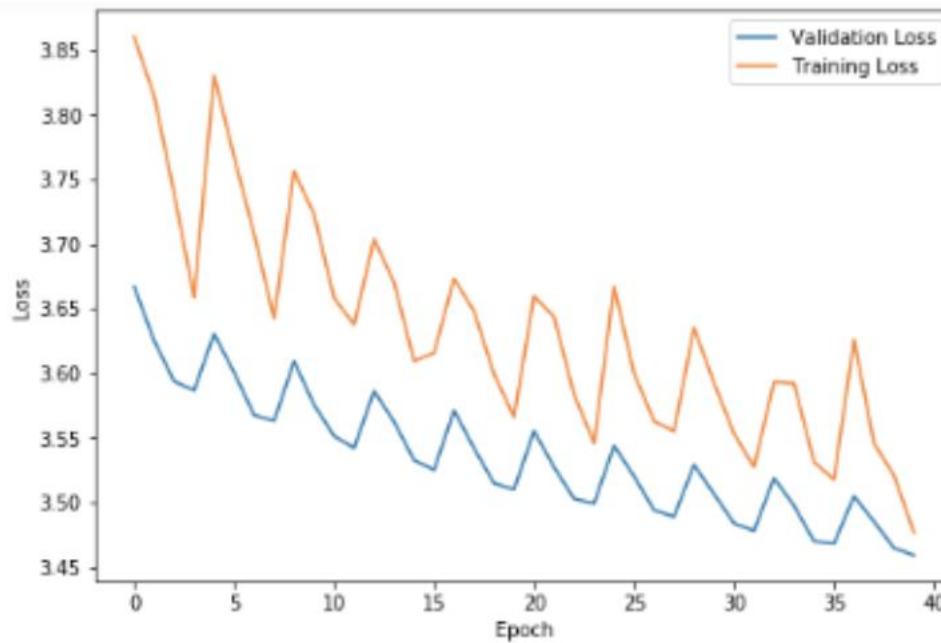
# Validation loss tracks train loss then diverges

- Q: overfitting or underfitting?
- A: overfitting from epoch 100 onwards
  - validation loss tracks training loss, but diverges from epoch 100
- solutions:
  - early stop at epoch 100
  - add more regularization to improve validation loss



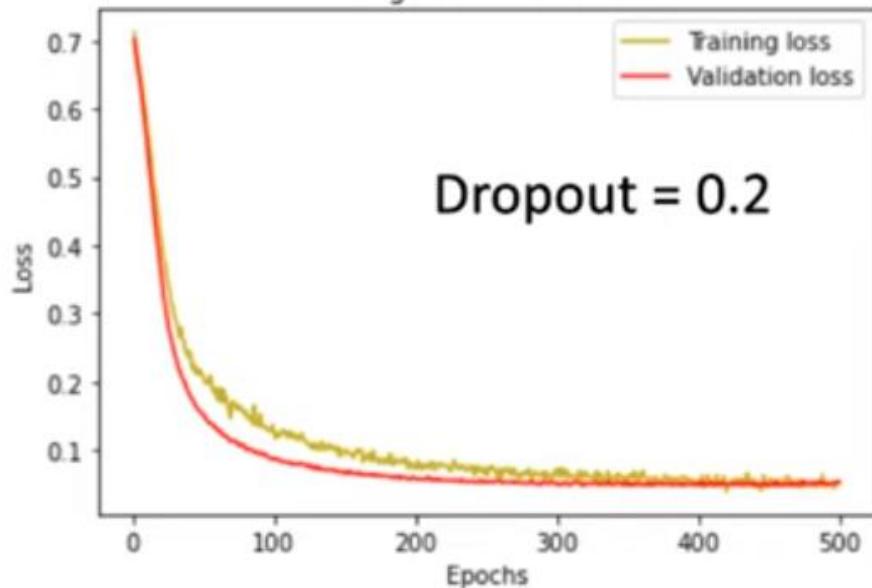
source: <https://www.youtube.com/watch?v=p3CcfIjycBA>

# Unusual train-validation loss curves



# Validation loss < training loss

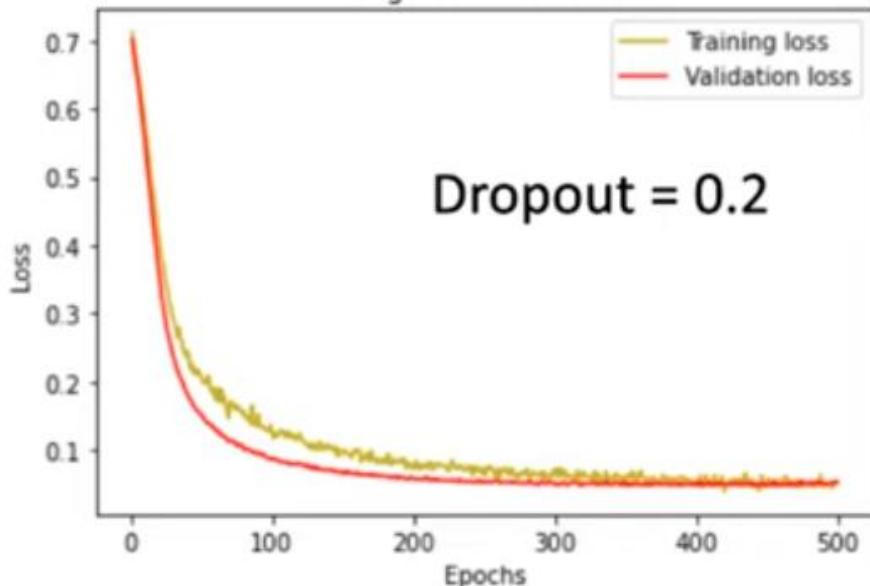
- Q: validation loss should normally be slightly higher than training loss, why is it lower here?



source: <https://www.youtube.com/watch?v=p3CcfIjycBA>

# Validation loss < training loss

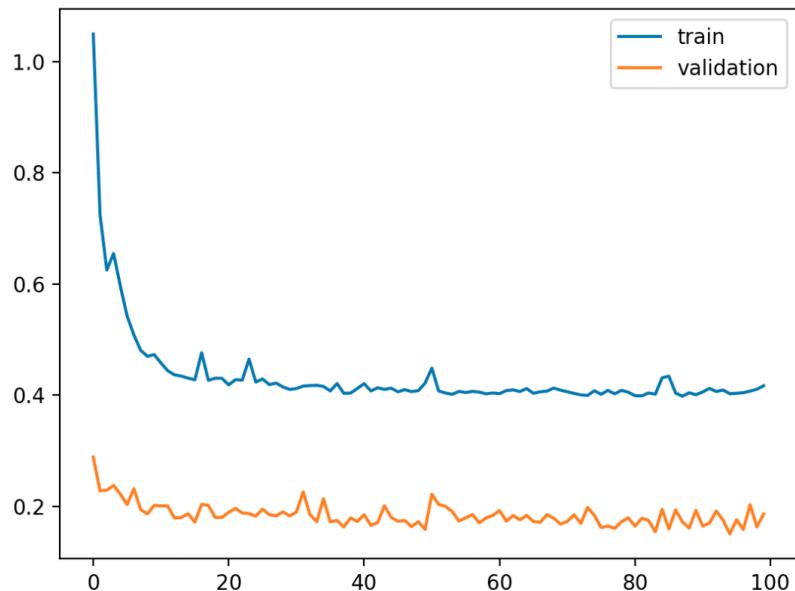
- Q: validation loss should normally be slightly higher than training loss, why is it lower here?
- A: the model is still a good fit if validation loss slightly lower than training loss
  - this effect may be due to Dropout, which disables some of the neurons during training
  - this lowers training performance at the expense of increased generalization during validation (when dropout is disabled)



source: <https://www.youtube.com/watch?v=p3CcfIjycBA>

# Validation loss << training loss

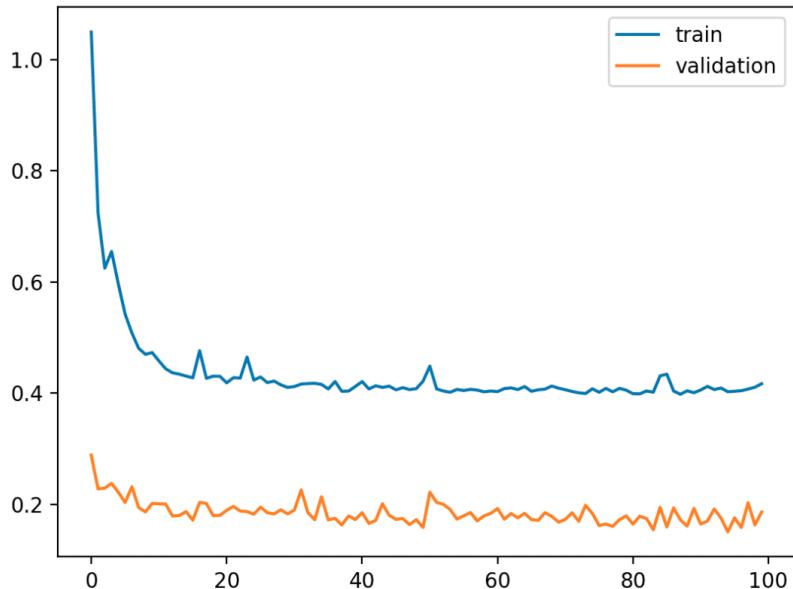
- Q: overfitting or underfitting?



source: <https://machinelearningmastery.com/learning-curves-for-diagnosing-machine-learning-model-performance/>

# Validation loss << training loss

- Q: overfitting or underfitting?
- A: neither
  - validation loss is much lower than training loss, which means the validation set is ‘easier’ for the model than the training set
  - this may point to **dataset issues**:
    - validation examples easier than training examples (due to some bug in the splitting process or augmenting training examples only)

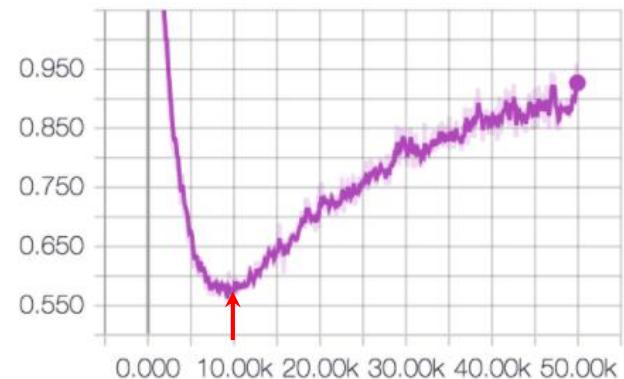


source: <https://machinelearningmastery.com/learning-curves-for-diagnosing-machine-learning-model-performance/>

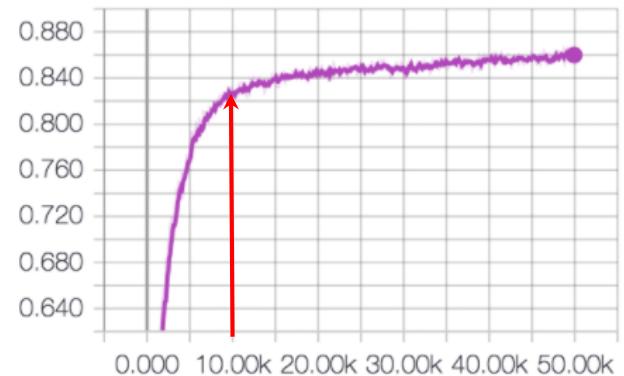
# Both validation loss & accuracy increase

- Q: lowest validation loss is at ~10k steps, but validation accuracy continues to increase after that, where does overfitting begin?

test/loss\_test



test/accuracy\_test

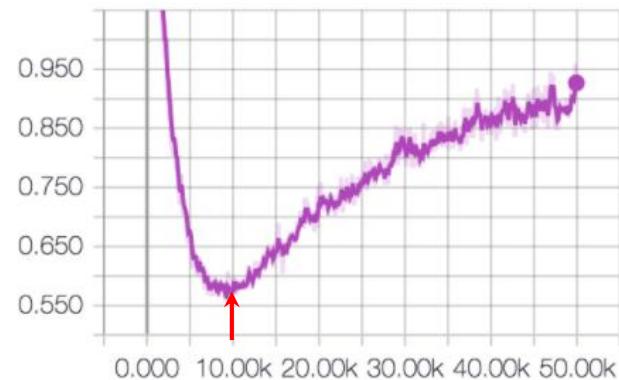


source

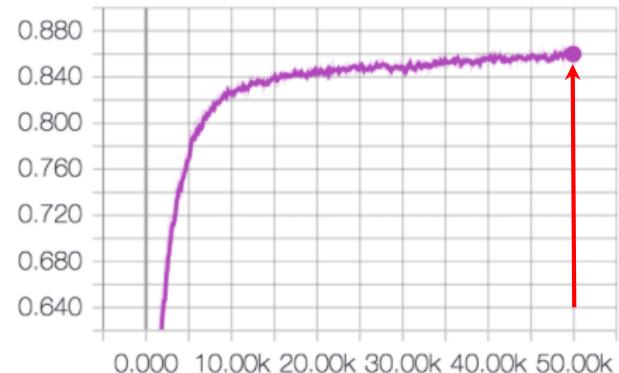
# Both validation loss & accuracy increase

- Q: lowest validation loss is at ~10k steps, but validation accuracy continues to increase after that, where does overfitting begin?
- A: overfitting begins at 10k steps
  - validation loss degrades rapidly, even though validation accuracy slightly improves

test/loss\_test



test/accuracy\_test

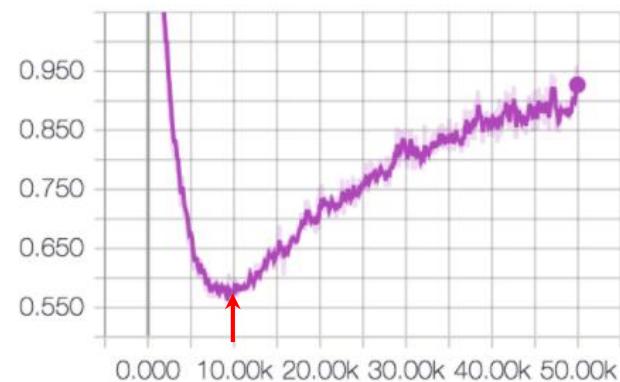


source

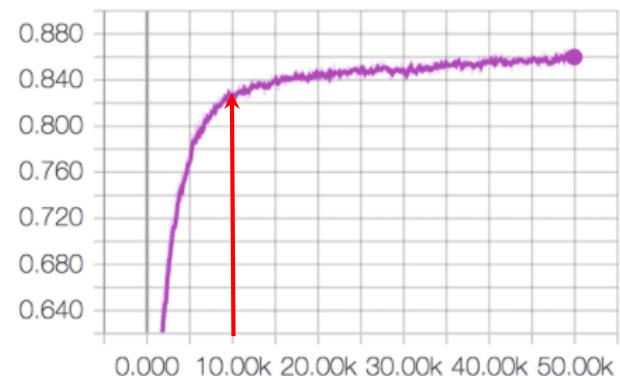
# Both validation loss & accuracy increase - why?

- loss is sensitive to probabilities:
  - it may become overconfident in wrong predictions, leading to higher loss
- thresholded metrics (like accuracy) ignore confidence
  - they usually decide right or wrong based on a fixed threshold (0.5)
  - accuracy may remain the same (or slightly improve) while probabilities deteriorate

test/loss\_test



test/accuracy\_test



source

# Both validation loss & accuracy increase - why?

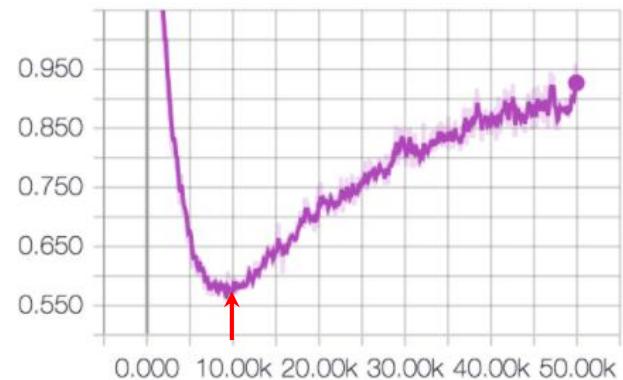
- becomes more confident in incorrect predictions ( $p=0.6 \rightarrow p=0.9$ )

	$p(y=0)$	$p(y=1)$	loss (-logp)
old predictions	0.6	0.4	0.92
new predictions	0.9	0.1	2.30

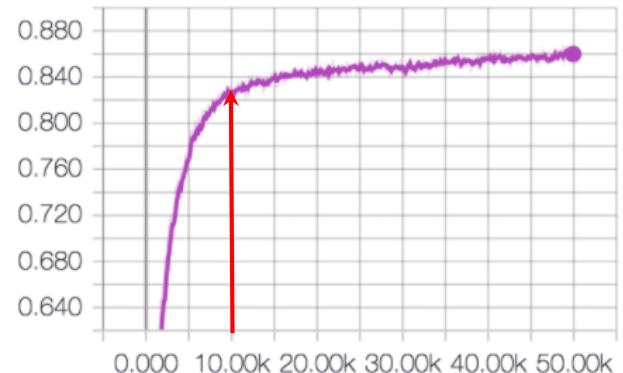
- wrong predictions close to threshold 0.5 become correct  $p=0.48 \Rightarrow p=0.51$

	$p(y=0)$	$p(y=1)$	loss (-logp)
old predictions	0.52	0.48	0.73
new predictions	0.49	0.51	0.67

test/loss\_test



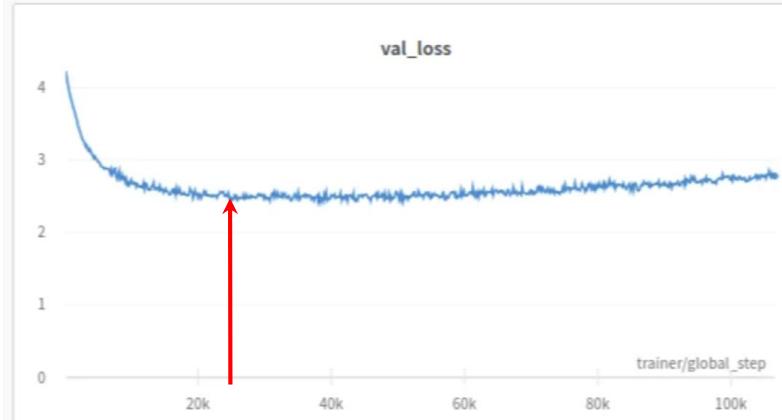
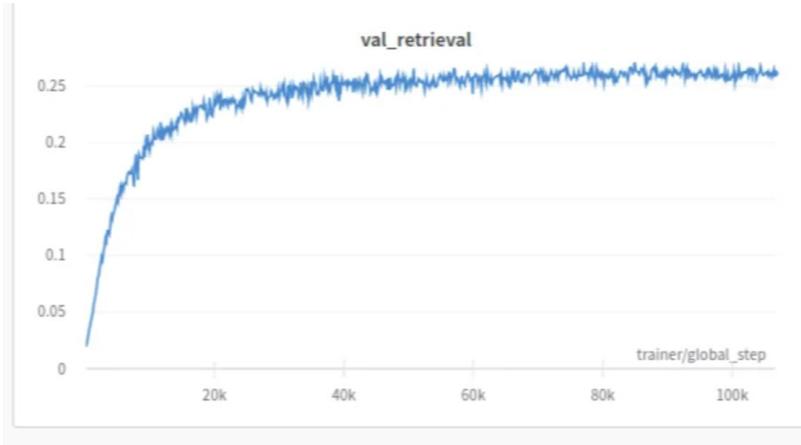
test/accuracy\_test



source

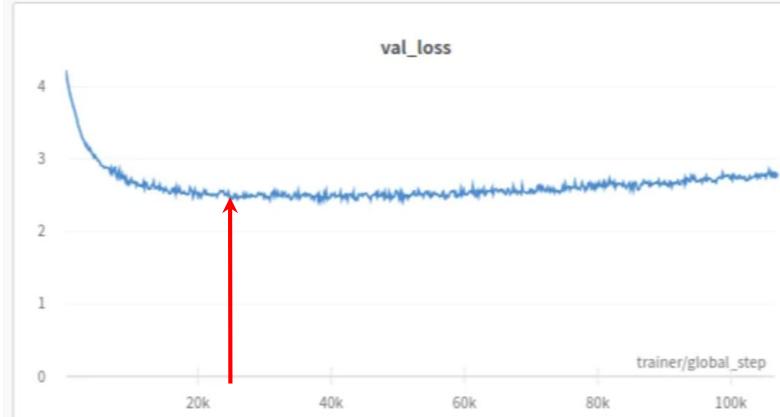
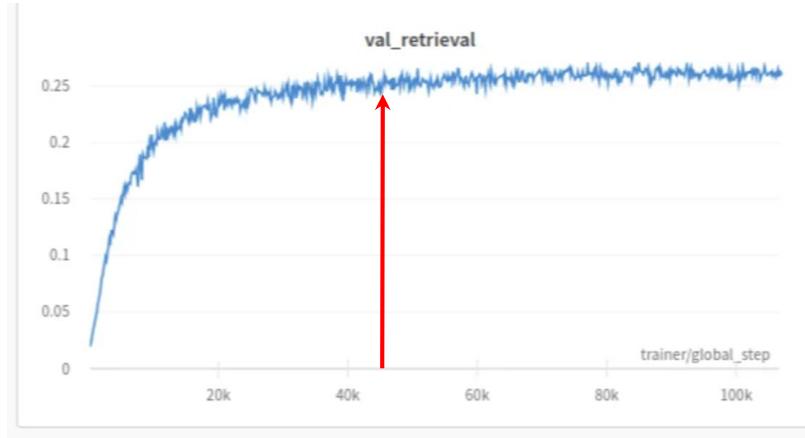
# Both validation loss & accuracy increase

- model begins to overfit around step 25K, which model to choose?



# Both validation loss & accuracy increase

- model begins to overfit around step 25K, which model to choose?
- we can pick a model with a slightly higher validation metric, as long as validation loss doesn't degrade significantly



# Learning curves TL;DR

- validate that the model is training correctly and generalizes:
  - **good fit:** train/validation loss decrease and stabilize at similar value
  - **beginning of overfit:** gap between validation and train loss widens

# Learning curves TL;DR

- validate that the model is training correctly and generalizes:
  - **good fit:** train/validation loss decrease and stabilize at similar value
  - **beginning of overfit:** gap between validation and train loss widens
- choose the best model:
  - **primary criterion:** choose epoch/step with *lowest loss on the validation set*

# Learning curves TL;DR

- validate that the model is training correctly and generalizes:
  - **good fit:** train/validation loss decrease and stabilize at similar value
  - **beginning of overfit:** gap between validation and train loss widens
- choose the best model:
  - **primary criterion:** choose epoch/step with ***lowest loss on the validation set***
  - **secondary check:** double check with validation metrics:
    - if *validation loss is at a minimum, or it slightly increases* and validation metric still increases, pick model where validation metric (acc./F1/etc.) maximum
    - if *validation loss rises significantly* and validation metric still increases, pick model where validation loss was minimum

# Debugging plan - recap

1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)

# Debugging plan - recap

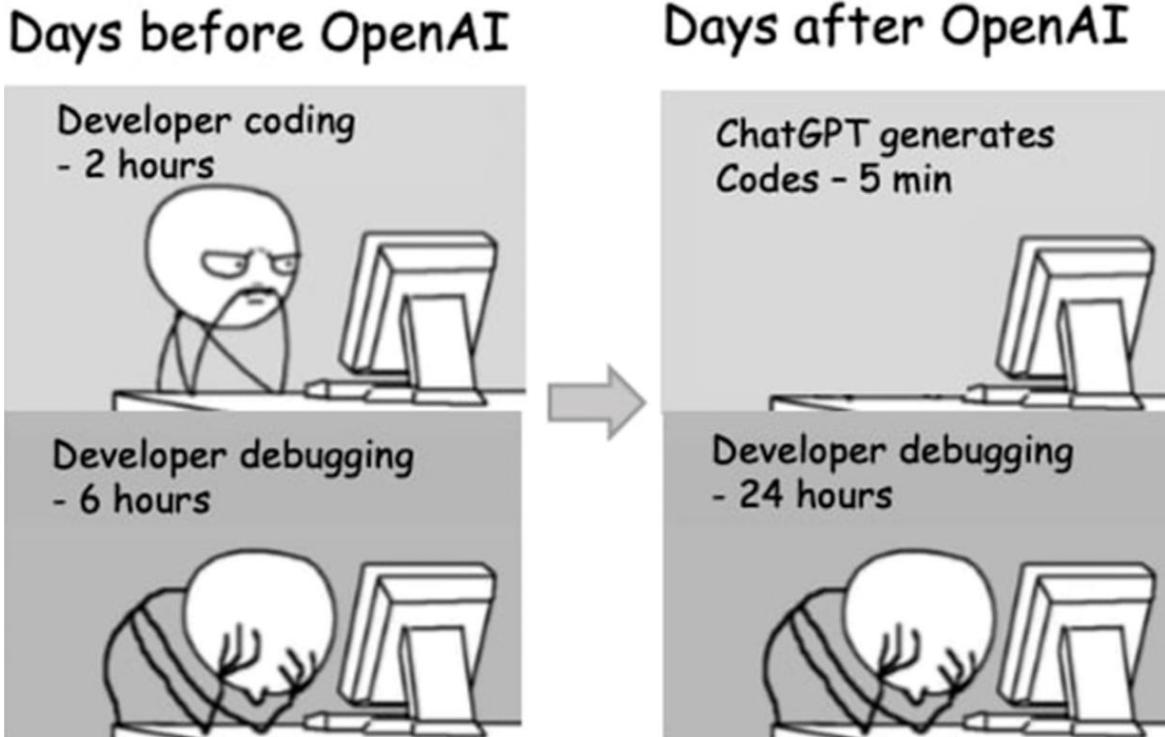
1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)
2. overfit model on a few (1-2) examples
  - o if train loss  $\rightarrow 0$ , start training on the full data
  - o otherwise look for **logical and dataset issues**

# Debugging plan - recap

1. run code and fix **syntactic bugs** (*shape/device/type mismatches, OOM issues*)
2. overfit model on a few (1-2) examples
  - o if train loss  $\rightarrow 0$ , start training on the full data
  - o otherwise look for **logical and dataset issues**
1. start training on the full dataset
  - o inspect loss/metric learning curves
  - o if loss curves don't show a good fit:
    - look for **hyperparameter misconfigurations** (learning rate, LR scheduler etc.)
    - reinvestigate **dataset issues** if model still can't generalize

# Debugging conclusion and QA

- ML projects have many sources of error:
  - start small and tackle one thing at a time 🔎
- embrace the debugging experience 🧐:
  - debugging is still a relevant skill even in the LLM era



<https://programmerhumor.io/debugging-memes/ai-revolution/>