

# Învățare Automată în Vedere Artificială

Curs 1: Introducere



# Echipă



Daniel Sociu



Alexandru Enache



Mihail Chirobocea



Andrei Janca



Vlad Păunescu

# Structură curs

- 9 săptamani (1 mai e miercuri, e in pre-sesiune)
- 8 cursuri x 2h
- 7 laboratoare x 2h (4 pct.)
  - 6 punctate, 1 bonus cu LLM-uri
  - se prezinta saptamana urmatoare (1 pct penalizare pt fiecare sapt)
- Proiect - 2 prezentări (6 pct.):
  - săptămâna 5 - prezentare temă proiect + SOTA + prezentare abordare aleasă 2.5 p
    - SOTA = State of the Art
  - săptămâna 10 - prezentare finală + demo 3.5 p
  - <https://paperswithcode.com/>
- Tehnologii: Python + PyTorch + Google Colaboratory

# Prerequisites

- Maths
  - Algebră Liniară
  - Analiză
  - Statistică și Probabilități
- Programming
  - Python
- Nice to have
  - Inteligență Artificială/Învățare Automată

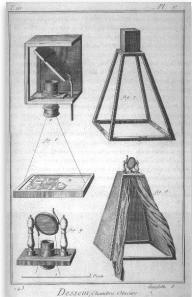


# Subiecte curs

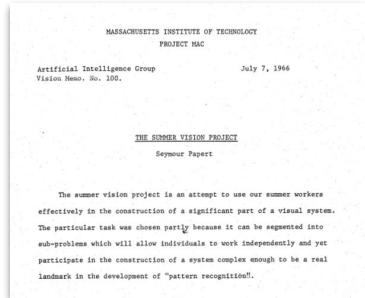
1. Introducere în Vederea Artificială
2. Clasificarea Imaginilor & Optimizare
3. Backpropagation & Convoluții
4. Rețele Convoluționale
5. Detectia de Obiecte
6. Segmentarea Semantică
7. Vision Transformers and GPT
8. Auto-encoders and diffusion models



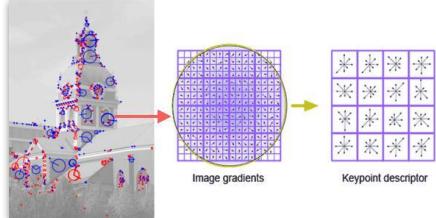
# Istoria Vederii (Artificiale) - Timeline



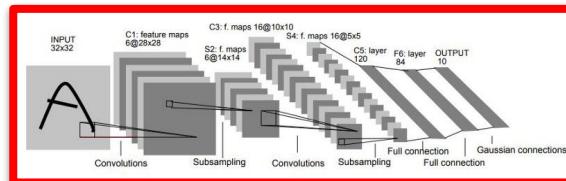
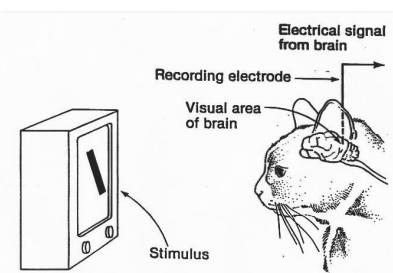
Camera Obscura



The Summer Vision Project

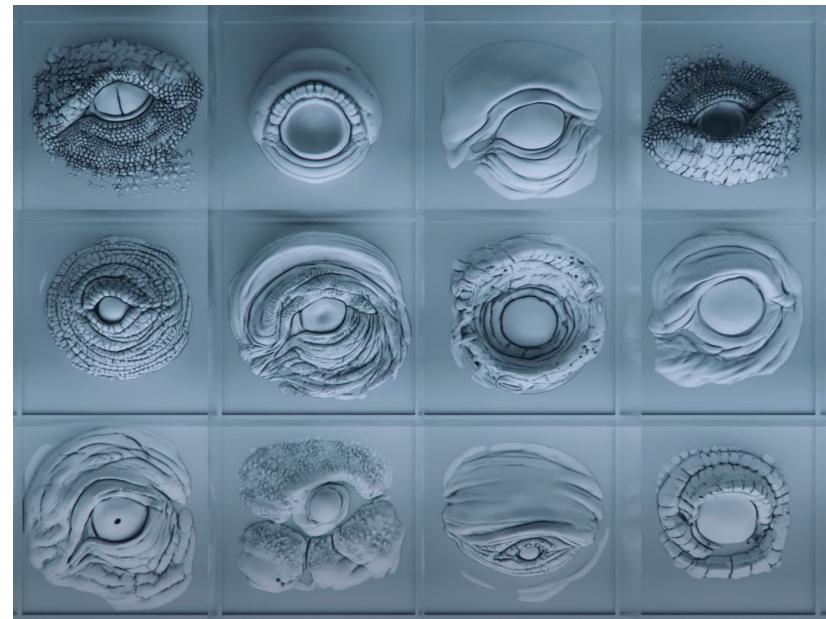
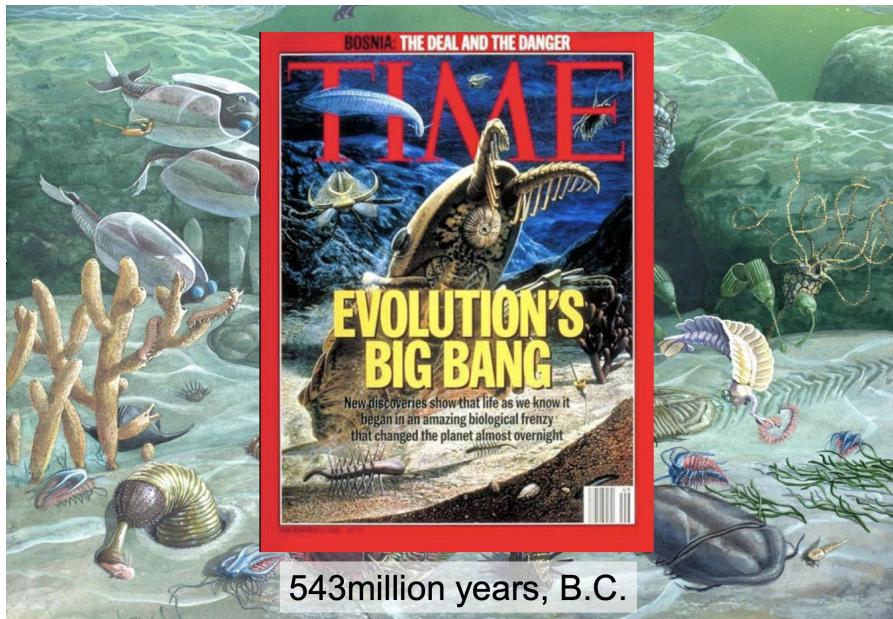


SIFT, David Lowe



# Istoria Vederii (Artificiale) - Explzia cambriană

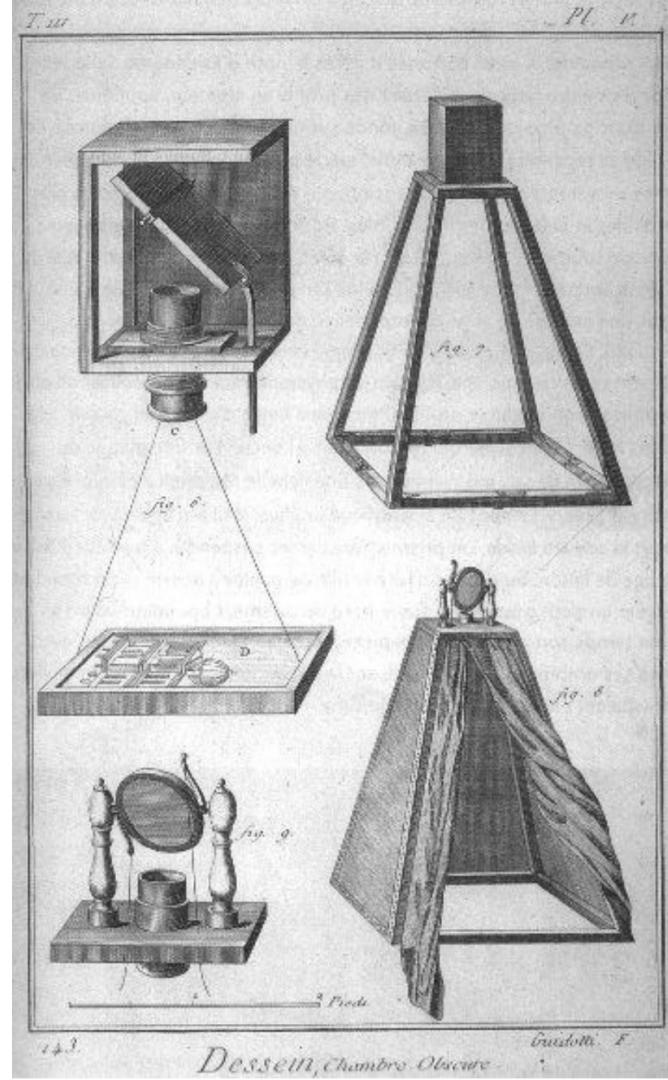
Eyesight have evolved multiple times, independently, in different species



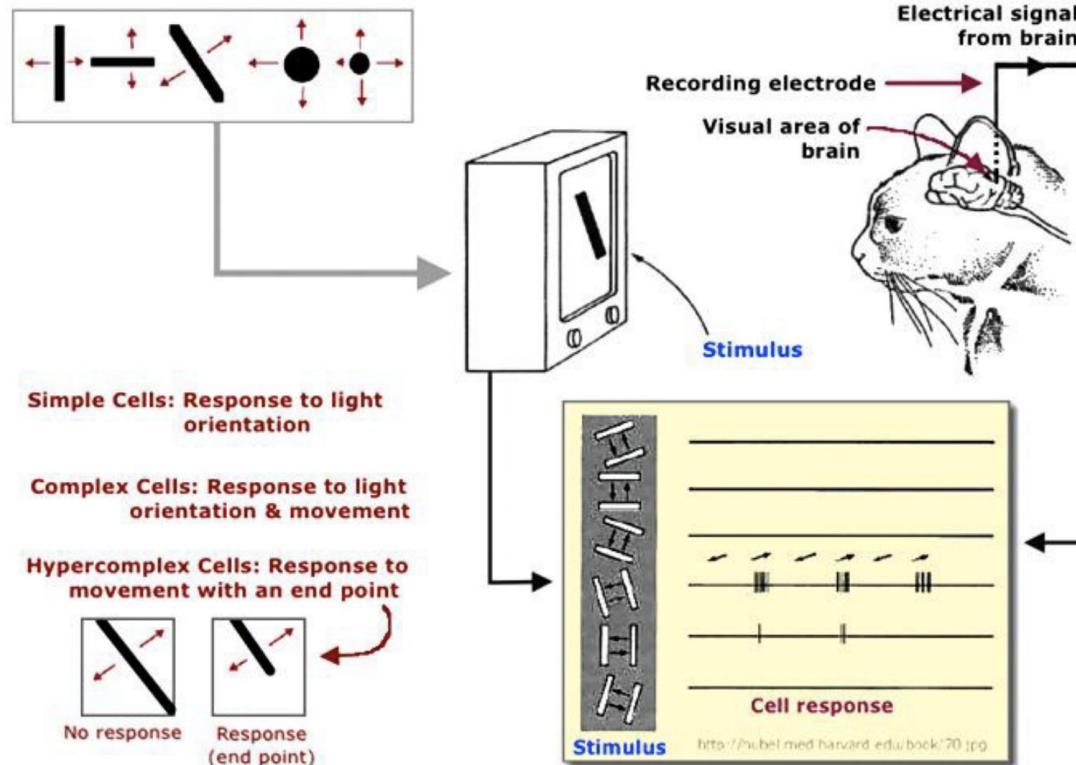
# Istoria Vederii (Artificiale)

## Camera Obscura

Leonardo da Vinci  
16th Century, A.D.



# Istoria Vederii (Artificiale) - Neuroștiințe

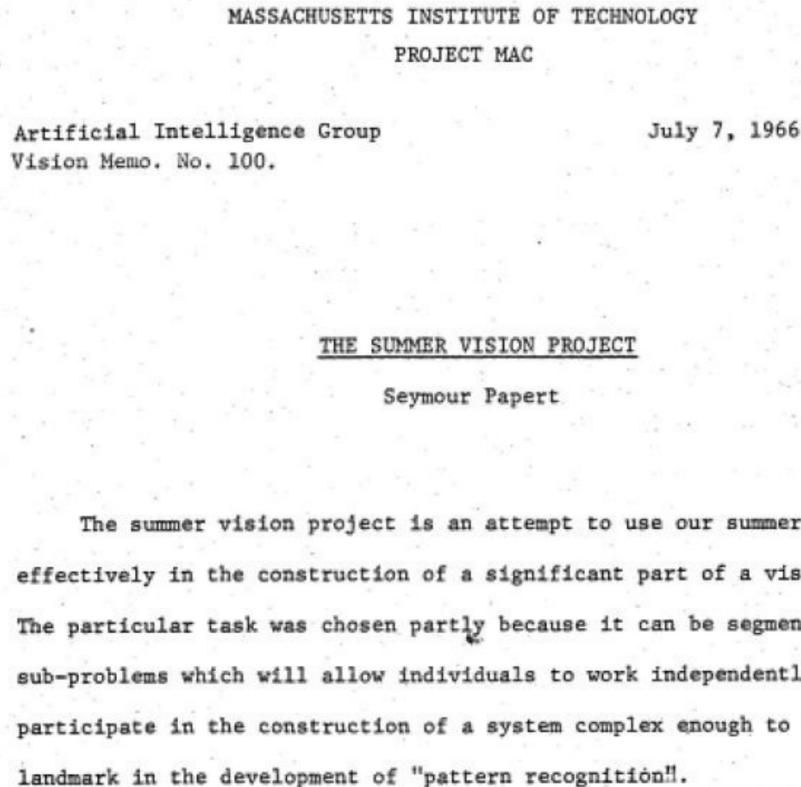


 's did it before it was cool

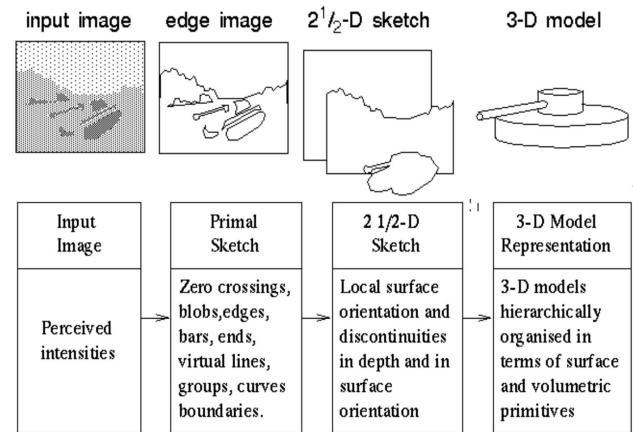
***Hubel & Wiesel, 1959***  
**video**

# Istoria Vederii (Artificiale)

## The Summer Vision Project MIT - 1966



## Vision - David Marr, 1970s



# Istoria Vederii (Artificiale) - Detectia fețelor

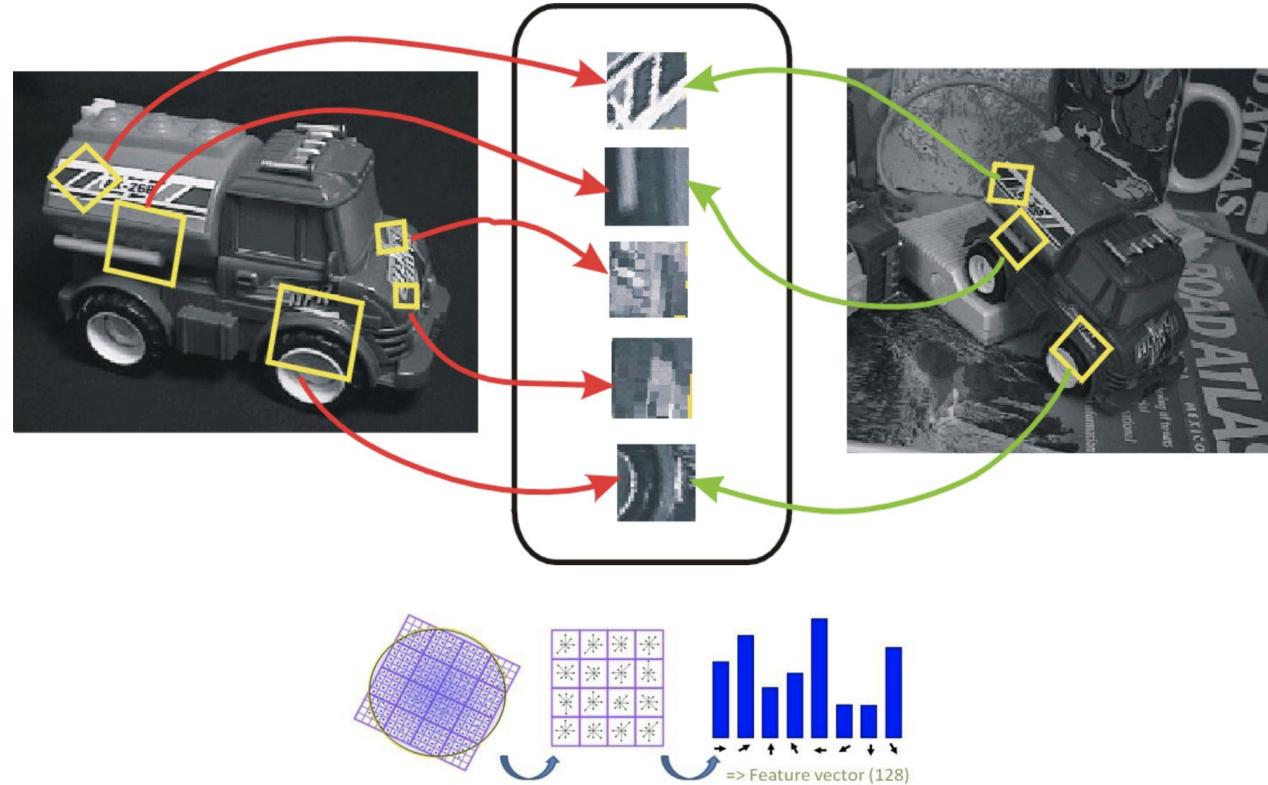
Face Detection,  
*Viola & Jones*  
**2001**



# Istoria Vederii (Artificiale) - Extragerea trăsăturilor

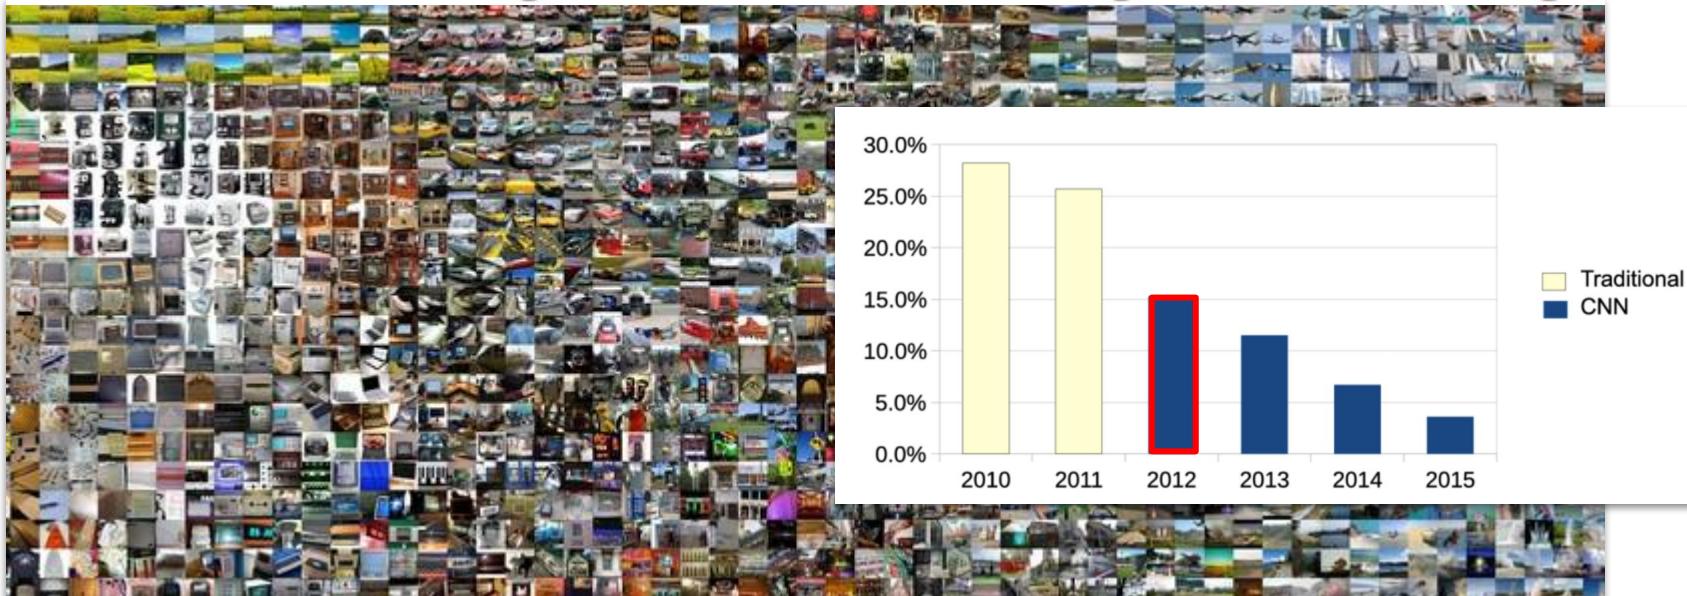
“SIFT” & Object  
Recognition,  
David Lowe, 1999

“Feature  
Extraction”



# Rețele Neurale Convoluționale

## IMAGENET Large Scale Visual Recognition Challenge

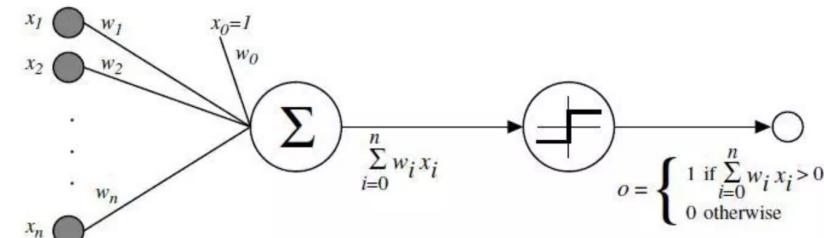


### The Image Classification Challenge

- 1.4 M Images
- 1,000 object classes

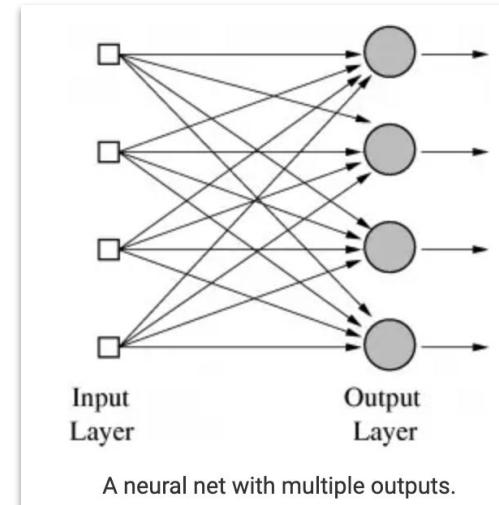
# Rețele Neurale

Frank Rosenblatt's **Perceptron**, 1958



'Artificial Neural Networks' - ANNs

- Extended to work for classification tasks with many categories
- Layers of Perceptrons - or neurons, or units, as they are usually called today



# Rețele Neurale Convolutionale

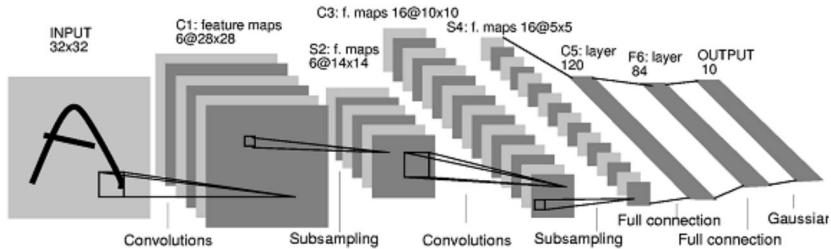
"According to the hierarchy model by Hubel and Wiesel, the neural network in the visual cortex has a hierarchy structure"

**Backpropagation algorithm -**  
Rumelhart, Hinton & Williams  
(1986)



1998

LeCun et al.



# of transistors



$10^6$

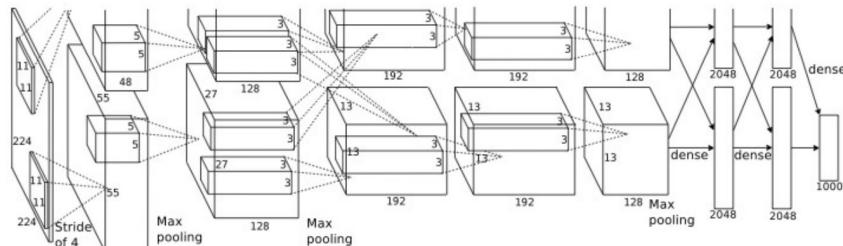
pentium® II

# of pixels used in training

$10^7$  NIST

2012

Krizhevsky  
et al.



# of transistors



$10^9$

GPUs



# of pixels used in training

$10^{14}$  IMAGENET

# Legea lui Moore in AI

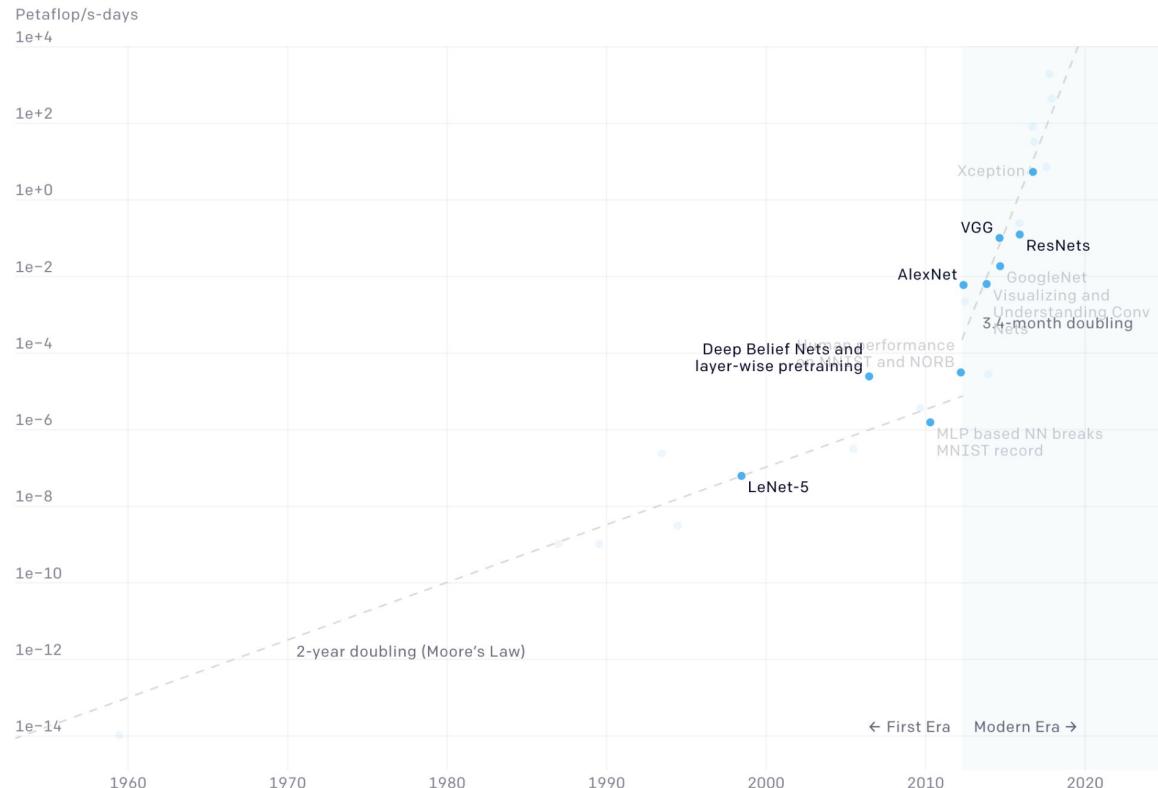
- the amount of compute used in the largest AI training runs has been increasing exponentially with a 3.4-month doubling time
- by comparison, Moore's Law had a 2-year doubling period
- Since 2012, this metric has grown by more than 300,000x
  - a 2-year doubling period would yield only a 7x increase
- Improvements in compute have been a key component of AI progress, so as long as this trend continues, it's worth preparing for the implications of systems far outside today's capabilities.

Source: <https://openai.com/blog/ai-and-compute/>



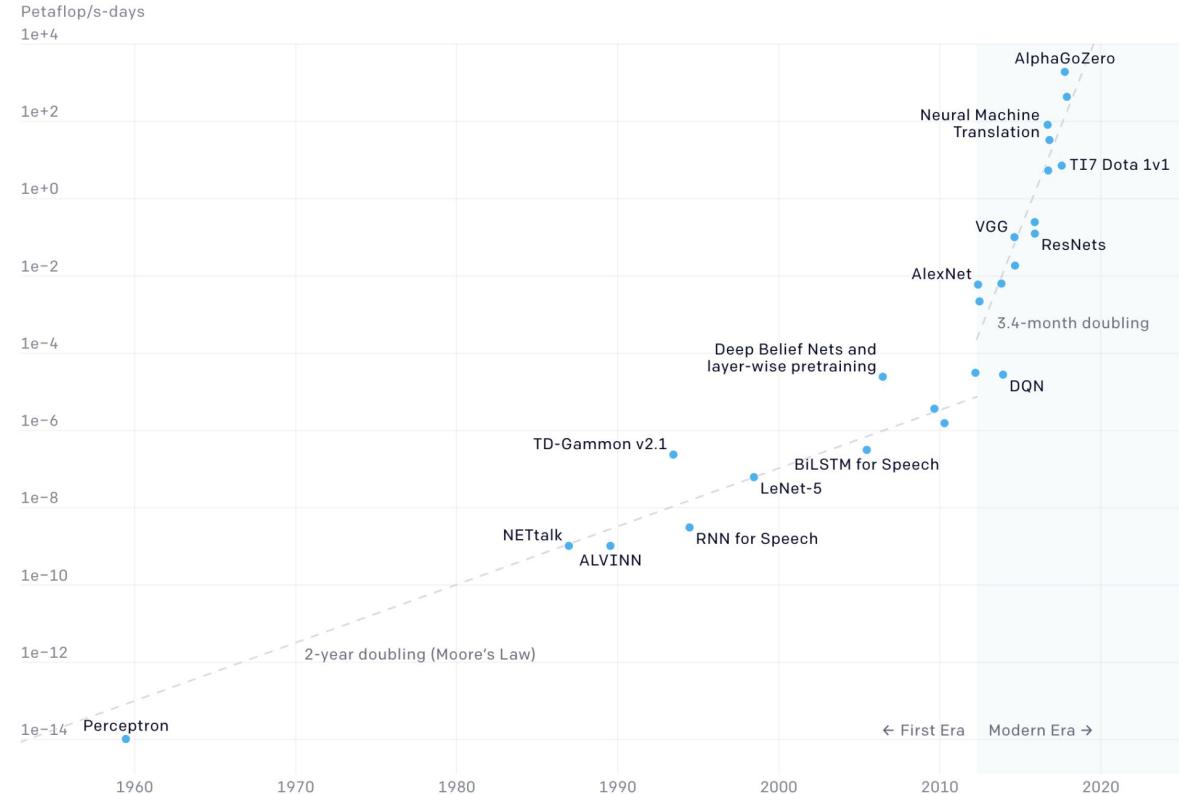
# Legea lui Moore in AI

Two Distinct Eras of Compute Usage in Training AI Systems (Vision)



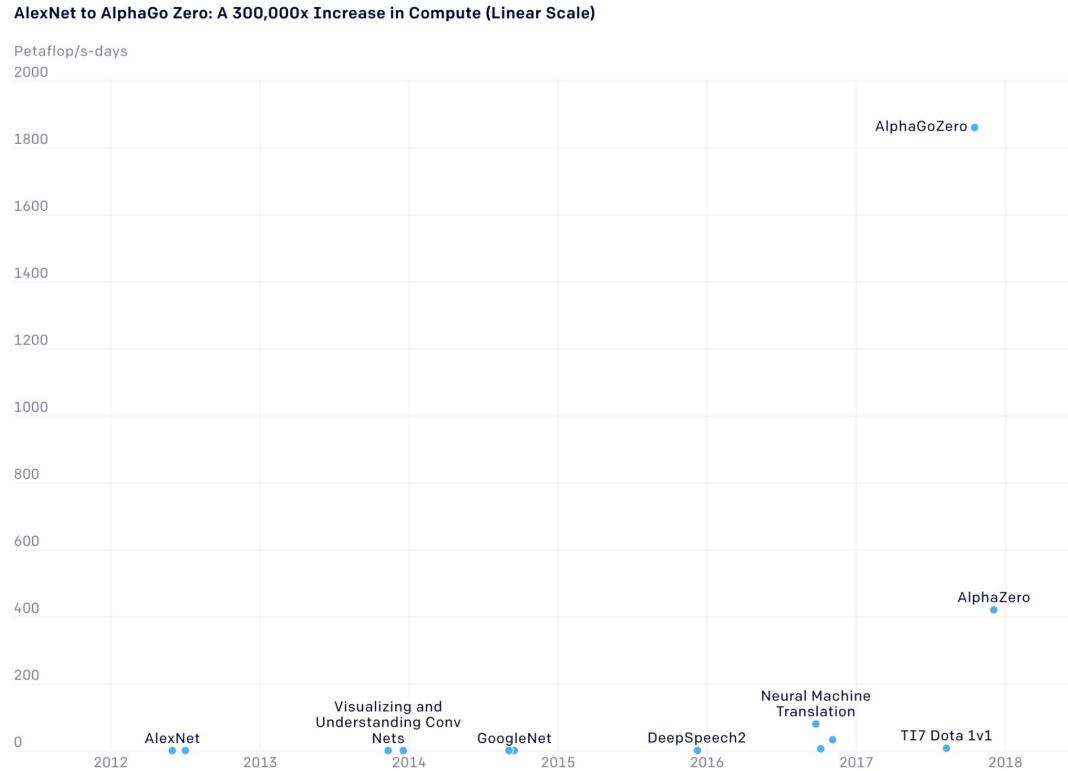
The total amount of compute, in petaflop/s-days,<sup>[2]</sup> used to train selected results that are relatively well known, used a lot of compute for their time, and gave enough information to estimate the compute used. Source <https://openai.com/blog/ai-and-compute/>

# Legea lui Moore in AI



The total amount of compute, in petaflop/s-days,<sup>[2]</sup> used to train selected results that are relatively well known, used a lot of compute for their time, and gave enough information to estimate the compute used. Source <https://openai.com/blog/ai-and-compute/>

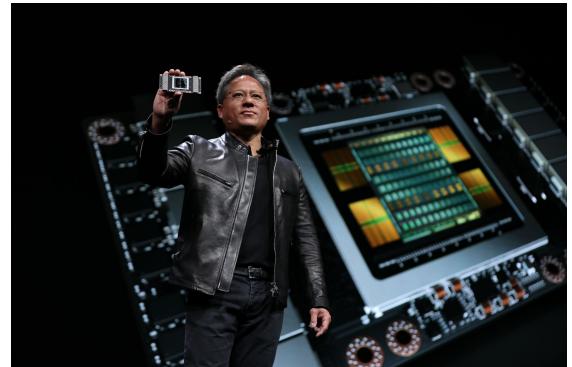
# Legea lui Moore in AI



The total amount of compute, in petaflop/s-days,<sup>[2]</sup> used to train selected results that are relatively well known, used a lot of compute for their time, and gave enough information to estimate the compute used. Source <https://openai.com/blog/ai-and-compute/>

# Modern Era of AI

- **How much compute to train the largest GPT-3(175B parameters) ?**
  - **3640 petaflop/s-day** - compute required to train largest GPT-3 model
  - **1 petaflop/s-day is equivalent to 8 Volta V100 GPUs at full efficiency of a day.**
    - **32GB x8 GPU Memory HBM2 [source](#)**
- How many days to train GPT-3.5?
  - As an example, it requires approximately 9.2 days on 512 V100 GPUs to train a 8.3B GPT-2 (Shoeybi et al., 2019),
  - 14.8 days on 10000 V100 GPUs to train a 175B GPT-3 (Patterson et al., 2021).
  - 3640 days on 8 V100 GPUs at full efficiency
  - OpenAI used Microsoft Azure cloud to train and scale



# Modern Era of AI

Mode	Parameters (B)	VRAM Requirement (GB)	Recommended GPU
DeepSeek-R1-Zero	671B	~1,543 GB	Multi-GPU setup (e.g., NVIDIA A100 80GB x16)
DeepSeek-R1	671B	~1,543 GB	Multi-GPU setup (e.g., NVIDIA A100 80GB x16)
DeepSeek-R1-Distill-Qwen-1.5B	1.5B	~3.9 GB	NVIDIA RTX 3060 12GB or higher
DeepSeek-R1-Distill-Qwen-7B	7B	~18 GB	NVIDIA RTX 4090 24GB or higher



# Aplicații

## Clasificare

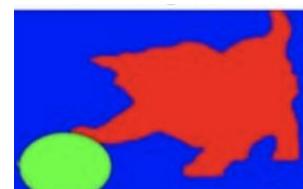


Cat

## Detectie



## Segmentare



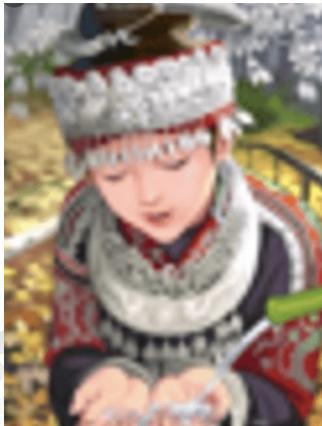
Red: cat  
Green: ball  
Blue: background

## Subtitrare



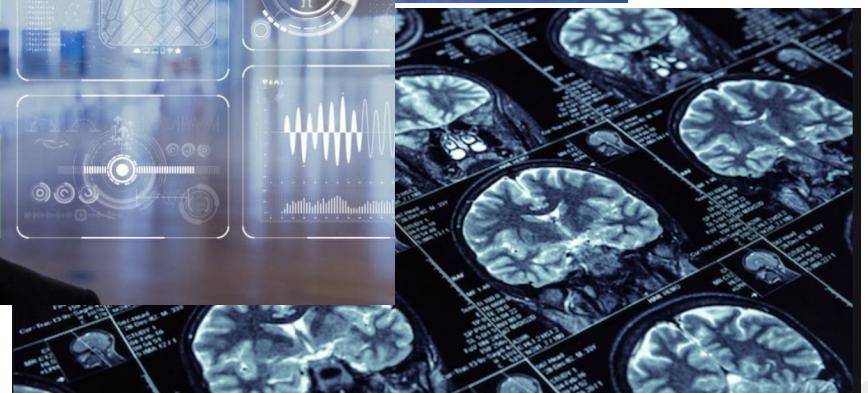
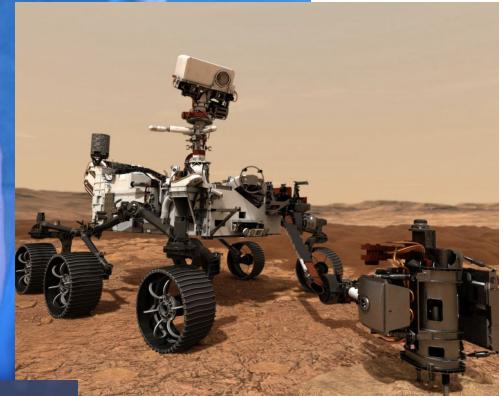
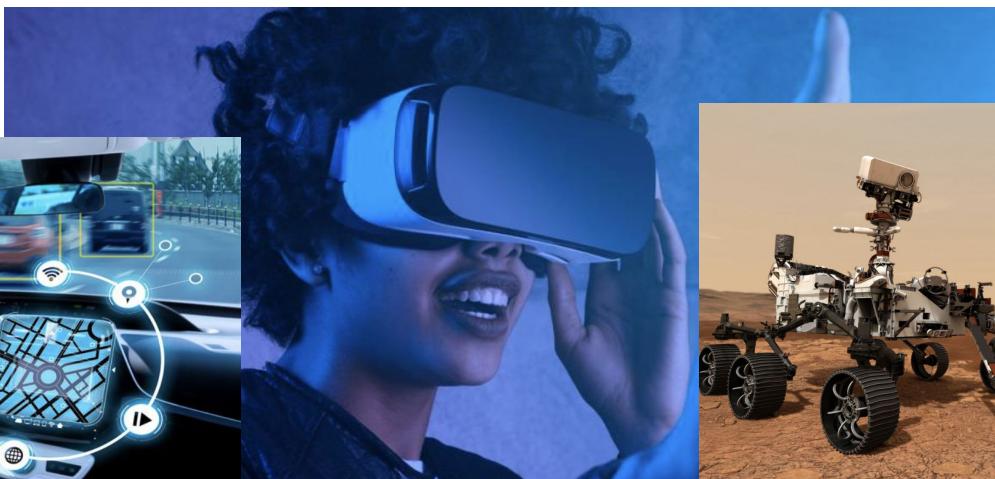
A cat is playing a ball.

## Super-rezoluție

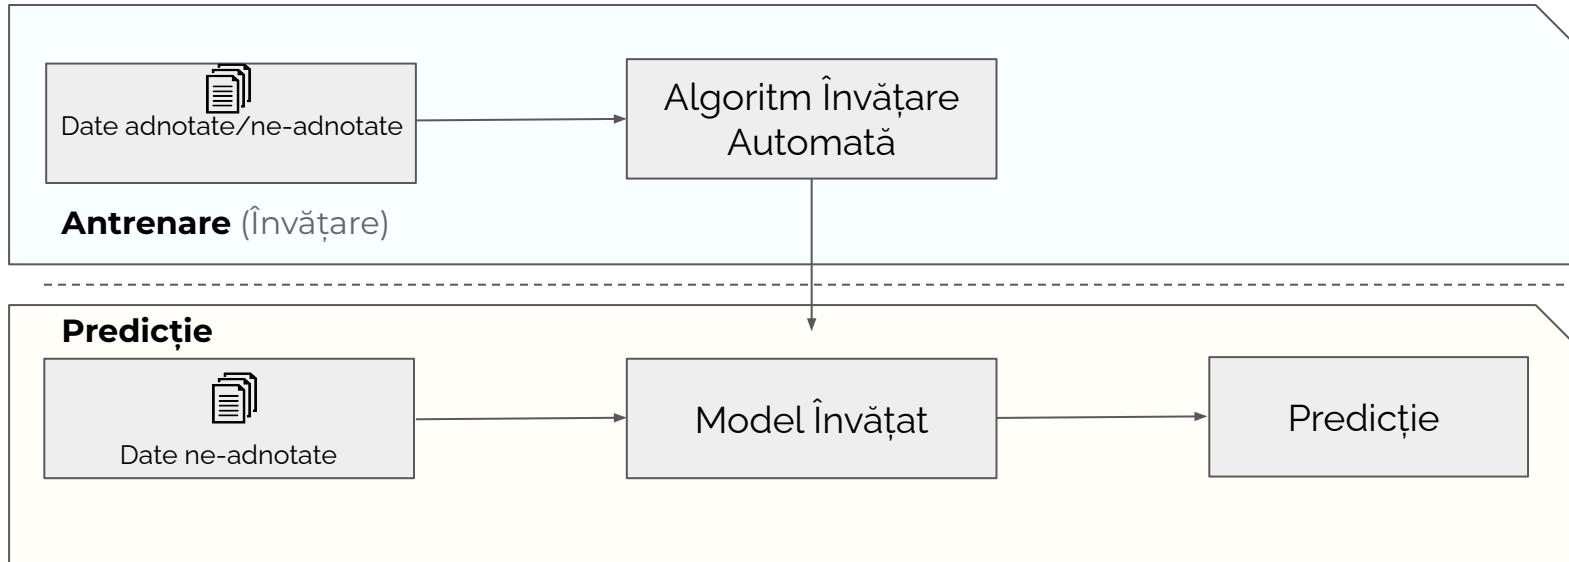


## Generare Imagini



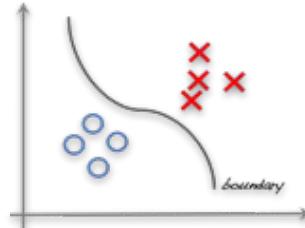


# Învățarea Automată - Concepte de Bază



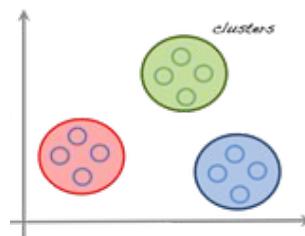
*"A computer program is said to learn from experience  $E$  with respect to some task  $T$  and some performance measure  $P$ , if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$ ."* - Tom Mitchell

# Învățarea Automată - Concepte de Bază



**Învățare supervizată (Supervised Learning):** Învățare dintr-un set de **date anotate**

Exemplu: detectare spam, estimarea prețului acțiunilor, recunoaștere de caractere



**Învățare nesupervizată (Unsupervised Learning):**  
Descoperirea tiparelor în **date neanotate**

Exemplu: segmentare clienți, detectare anomalii



**Învățare prin Recompensă (Reinforcement Learning):**

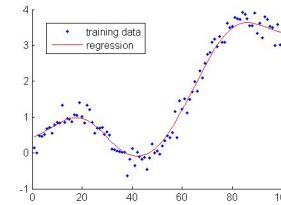
Exemplu: învăță să joace Tetris, manipulare roboți

# Învățarea Automată - Concepte de Bază

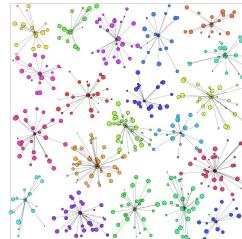


pisică 0.9  
câine 0.1

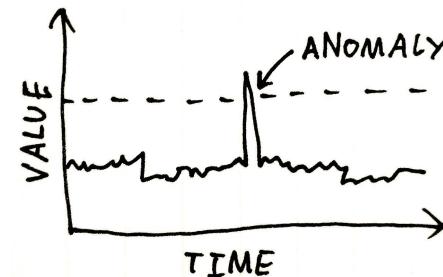
Clasificare  
(Supervizată - predictivă)



Regresie  
(Supervizată - predictivă)



Grupare (Clustering)  
(Nesupervizată - descriptivă)



Detectia anomaliilor  
(Nesupervizată - descriptivă)

# Regresia liniara

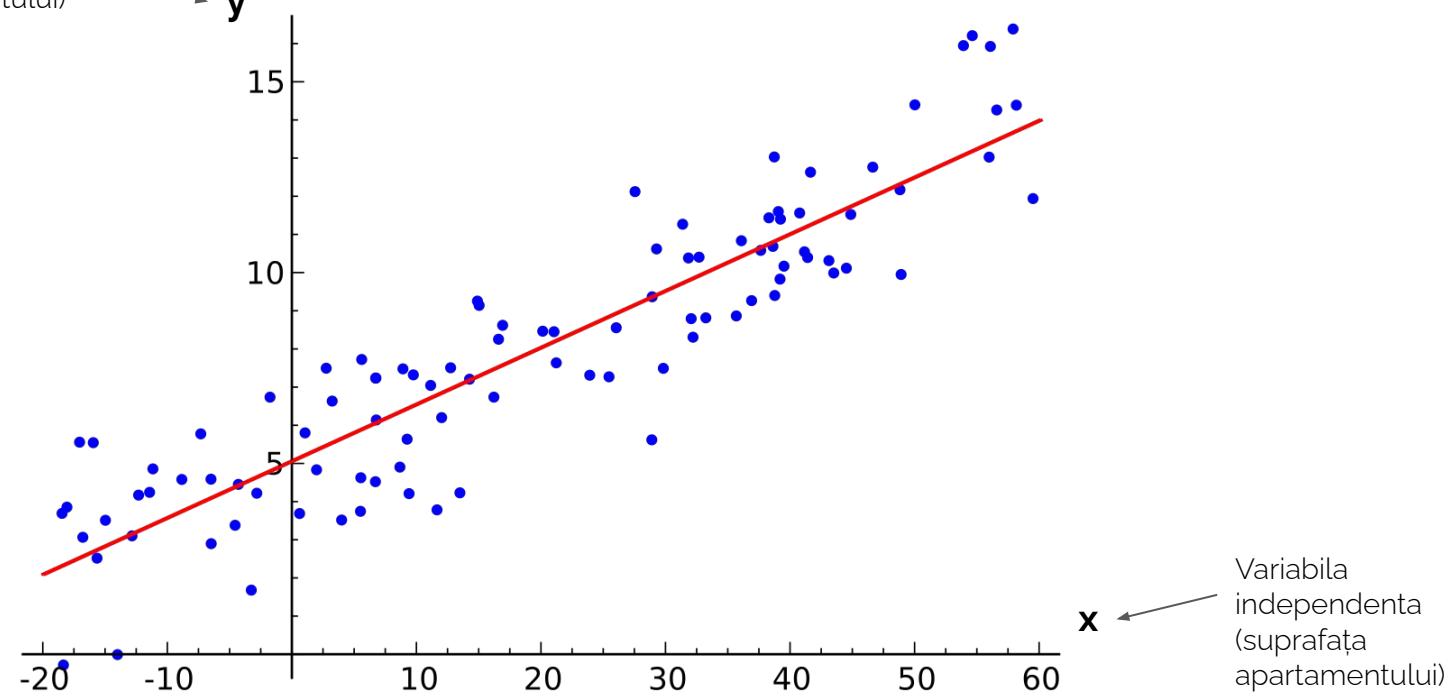
Regresia: prezicerea valorii unei variabile continue

Suprafata [mp]	Pret [* 10.000\$]
40	5
35	4.7
67	12
...	...
29	3.8

Q: Cum putem prezice prețul unui apartament a cărui suprafață nu este inclusă în setul de date?

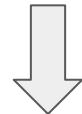
# Regresia liniara

Variabila dependenta  
(ex.: prețul apartamentului)

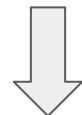


# Regresia liniara

Setul de antrenare: n instante de tipul {x; y}



Algoritm de învățare



Input

→ Functie care reflecta dependenta lui Y fata de X →

Predictie



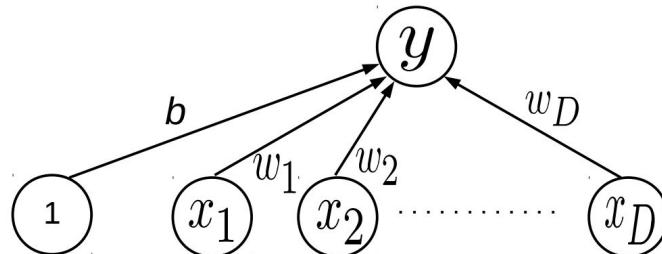
# Optimizare - Regresie Liniară

- **Modelele liniare** sunt componente ale bază ale rețelelor neuronale
- Vom formula învățarea ca o problemă de **optimizare** (cu privire la acești parametri)
- \*Putem adăuga o constantă '1' pentru fiecare intrare (feature) și scrie ca:

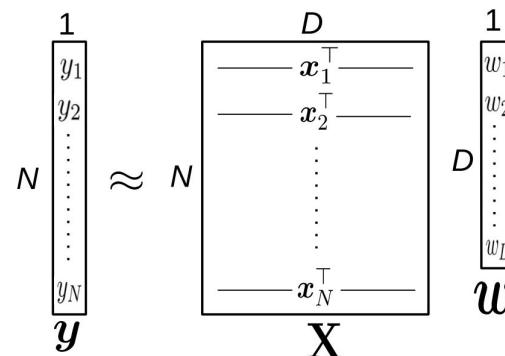
$$y = \mathbf{w}^\top \mathbf{x}, \quad \mathbf{x}, \mathbf{w} \in \mathbb{R}^{D+1}$$

- Pentru antrenare, presupunem existența unui set de **date de antrenare**

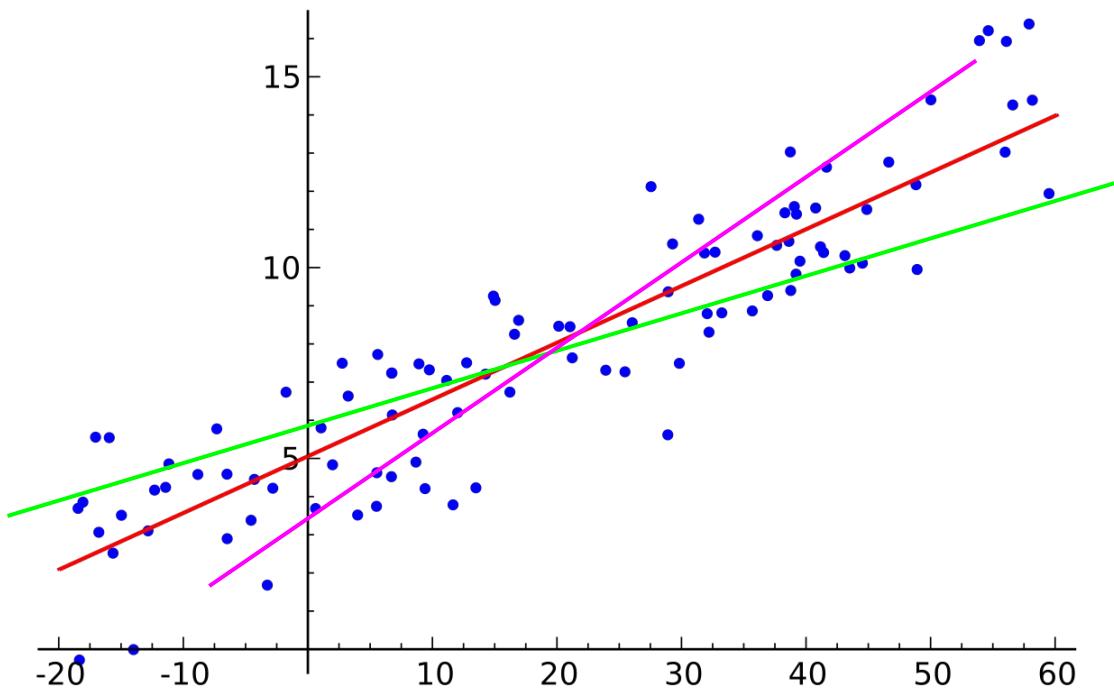
- $\{(\mathbf{x}_n, y_n)\}_{n=1}^N \quad \mathbf{x}_n \in \mathbb{R}^D, \quad y_n \in \mathbb{R}$



$$y = \sum_{d=1}^D w_d x_d + b = \mathbf{w}^\top \mathbf{x} + b$$



# Regresia liniara



# Regresia Liniară

- Eroarea modelului (reziduu) pentru un exemplu:

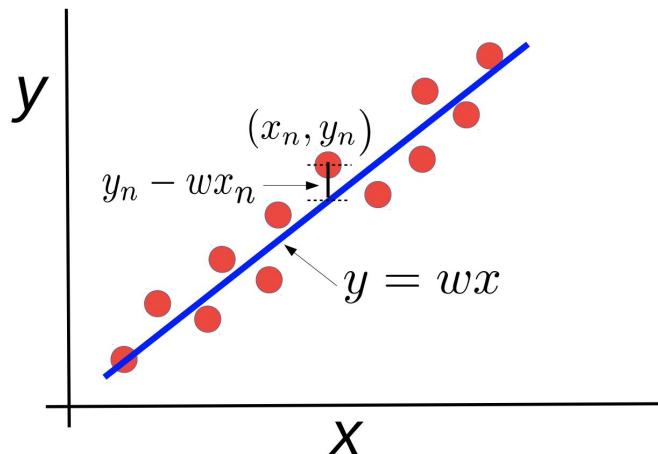
$$y_n - \mathbf{w}^\top \mathbf{x}_n \quad (\text{sau } y_n - \mathbf{w}\mathbf{x}_n \text{ pentru un input scalar - Fig})$$

- Definim eroarea totală sau „pierdere” pe întreg setul de antrenare (**suma reziduurilor**):

$$\circ \quad \mathcal{L}(\mathbf{w}) = \sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2$$

- Optimizarea modelului** constă în găsirea celui mai bun  $\mathbf{w}$ - care minimizează eroarea de mai sus

$$\circ \quad \hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \arg \min_{\mathbf{w}} \sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2$$



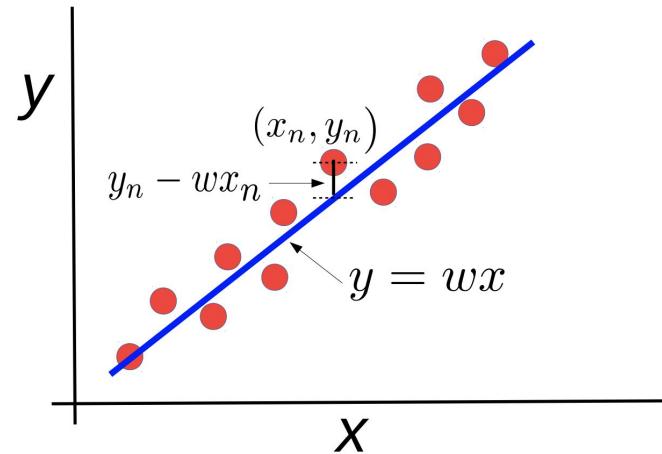
# Metoda celor mai mici pătrate (“Least Squares”)

- Abordare standard în analiza de regresie pentru a aproxima soluția sistemelor nedeterminate
  - Seturi de ecuații în care există mai multe ecuații decât necunoscute
  - Minimizarea sumei pătratelor reziduurilor realizate în rezultatele fiecărei ecuații

- Luând derivata (gradientul) lui  $L(w)$  în raport cu  $\mathbf{w}$  și setand la zero (pentru a găsi valoarea minimă)
  - $\sum_{n=1}^N 2(y_n - \mathbf{w}^\top \mathbf{x}_n) \frac{\partial}{\partial \mathbf{w}} (y_n - \mathbf{x}_n^\top \mathbf{w}) = 0 \Rightarrow \sum_{n=1}^N \mathbf{x}_n (y_n - \mathbf{x}_n^\top \mathbf{w}) = 0$

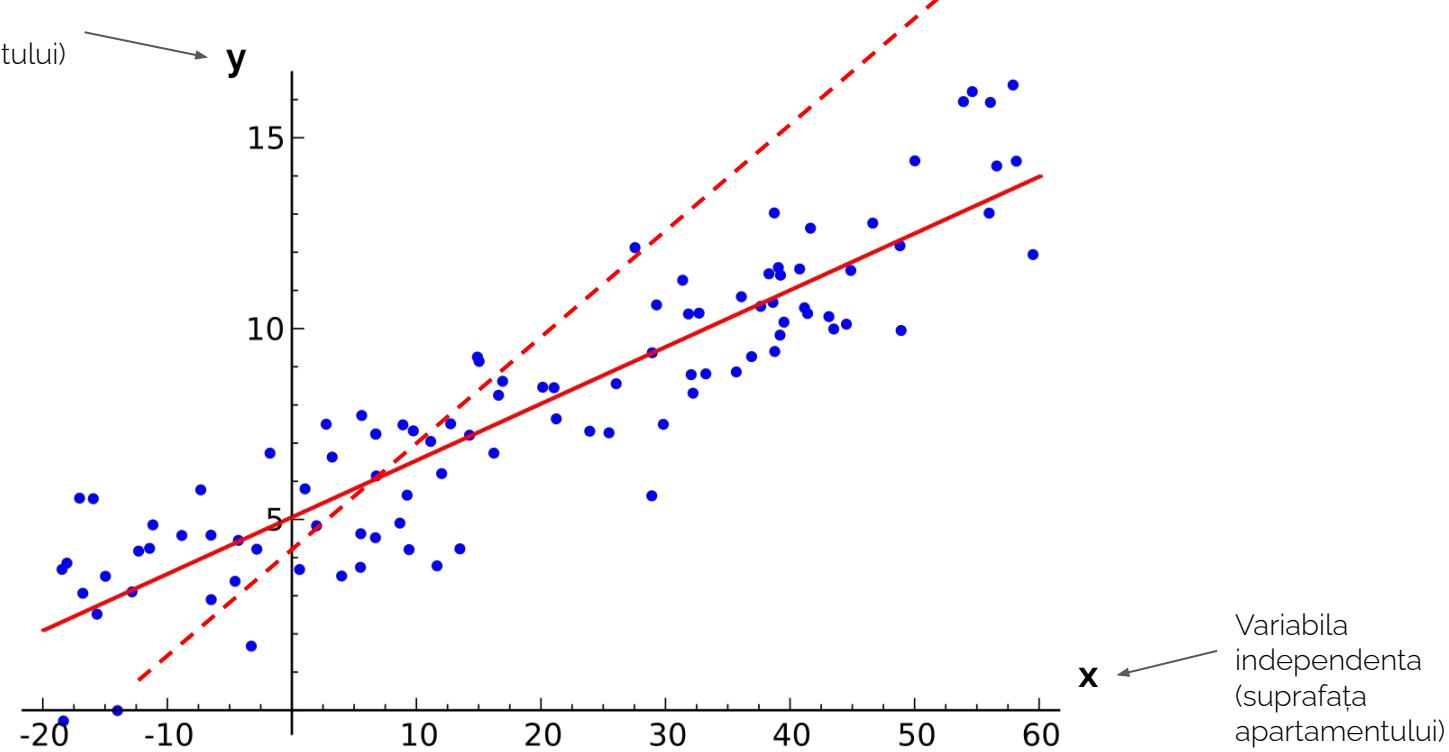
- Simplificând în continuare, obținem o soluție de formă închisă pentru  $\mathbf{w}$ 
  - $\mathbf{w} = (\sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^\top)^{-1} \sum_{n=1}^N y_n \mathbf{x}_n = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$

$$\mathcal{L}(\mathbf{w}) = \sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2$$



# Regresia liniara

Variabila dependenta  
(ex.: prețul apartamentului)



# Gradient Descent

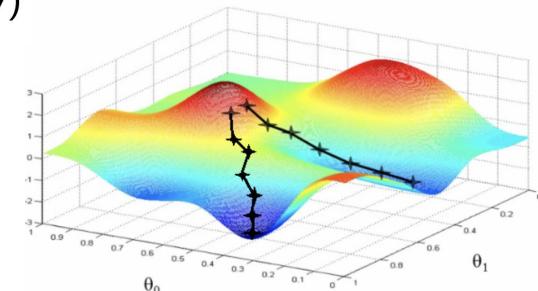
- Soluția de formă închisă are următoarele probleme:
 
$$(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$
  - Matricea  $(\mathbf{X}^T \mathbf{X})$  s-ar putea să nu fie inversabilă
  - Bazându-ne exclusiv pe minimizarea erorii de antrenament putem face **"overfit"** pe setul de antrenare
  - Poate fi foarte ineficientă din punct de vedere computațional atunci când D (dimensionalitatea) este foarte mare
- O modalitate mai rapidă și eficientă este de a utiliza optimizarea iterativă - gradient stocastic (**"Gradient Descent"**).

- Se porneste cu o valoare inițială (aleatoare)  $\mathbf{w} = \mathbf{w}^{(0)}$
- Actualizăm W deplasându-ne de-a lungul gradientului funcției de pierdere  $\mathcal{L}$

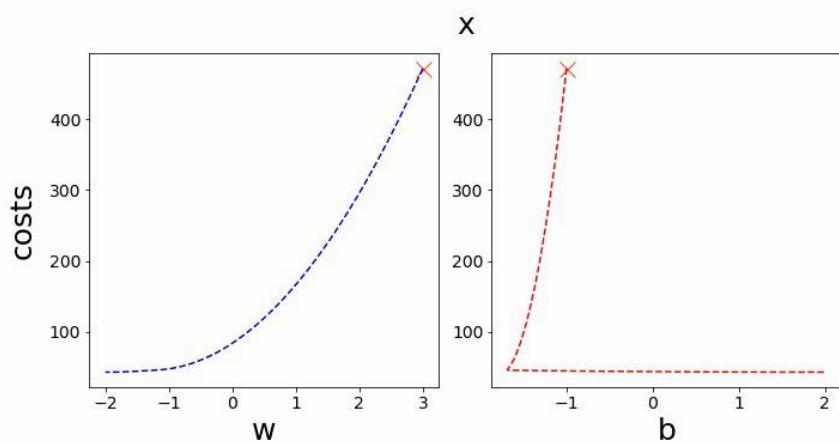
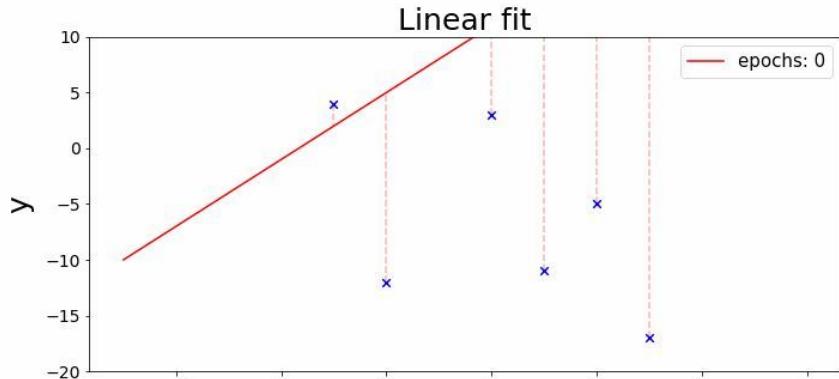
$$\boxed{\mathbf{w}^{(t)} = \mathbf{w}^{(t-1)} - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{w}} \Big|_{\mathbf{w}=\mathbf{w}^{(t-1)}}}$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = - \sum_{n=1}^N \mathbf{x}_n (y_n - \mathbf{x}_n^T \mathbf{w})$$

- Unde,  $\eta$  este rata de învățare (pasul fiecărei iterări)
- Repetăm până la convergență



# Gradient Descent



Credits  
Gradient descent animation

# Rezumat

- Loss function
- Optimizare - antrenare (găsirea parametrilor)
- Least Squares
- Gradient Descent



# Resurse

- Deep Learning course : [CS230 Deep Learning](#)
- Deep Learning course: [CS231n: Convolutional Neural Networks for Visual Recognition](#)
- Deep Learning book : <https://www.deeplearningbook.org>
- Mathematics book : <https://mml-book.github.io/>
- Machine Learning book (advanced) : [Pattern Recognition And Machine Learning - Bishop, 2006](#)
- But what is a GPT? <https://www.3blue1brown.com/lessons/gpt>
- 

**Seek out guidance and further resources from the team :)**

